



# Computer Vision



**Lecture 4: Machine Learning**  
Oliver Bimber

# Last Week: Digital Image Processing

## Convolution



**JKU** JOHANNES KEPLER  
UNIVERSITY LINZ

## Spatial vs. Gradient Domain

Image gradient:

$$\nabla I = \left( \frac{\partial I}{\partial x}, \frac{\partial I}{\partial y} \right)$$

Gradient Domain

Derivative / Integral

$$I = \int \nabla I$$

$$\begin{aligned} \text{div}(\nabla I) &= \frac{\partial I_x}{\partial x} + \frac{\partial I_y}{\partial y} = I_{xx} + I_{yy} = \nabla^2 I \\ \text{curl}(\nabla I) &= \frac{\partial I_y}{\partial x} - \frac{\partial I_x}{\partial y} = I_{yx} - I_{xy} \end{aligned}$$

Processing the Vectorfield is called  
Gradient Domain Processing.


Integrating this Vectorfield results in the original Image as long as the Vectorfield has zero Curl. In Practice: that does not work as the Vectorfields is no longer conservative anymore after processing it (it has non-zero Curl).

Instead: solve 2D Poisson Equation:  $\nabla^2 I = \text{div}(G)$

**JKU** JOHANNES KEPLER  
UNIVERSITY LINZ


**JKU** JOHANNES KEPLER  
UNIVERSITY LINZ

## Spatial vs. Frequency Domain



**JKU** JOHANNES KEPLER  
UNIVERSITY LINZ

## Convolutional Neural Networks (CNNs)




**JKU** JOHANNES KEPLER  
UNIVERSITY LINZ

# Course Overview

CW	Topic	Date	Place	Lab
41	Introduction and Course Overview	07.10.2025	Zoom	Lab 1
42	Capturing Digital Images	14.10.2025	Zoom	Lab 2
43	Digital Image Processing	21.10.2025	Zoom	Assignment 1
→ 44	Machine Learning	28.10.2025	Zoom	
45	Feature Extraction	04.11.2025	Zoom	Open Lab 1
46	Segmentation	11.11.2025	Zoom	Assignment 2
47	Optical Flow	18.11.2025	Zoom	Open Lab 2
48	Object Detection	25.11.2025	Zoom	Assignment 3
49	Multi-View Geometry	02.12.2025	Zoom	Open Lab 3
50	3D Vision	09.12.2025	Zoom	Assignment 4
3	Trends in Computer Vision	13.01.2026	Zoom	
4	Q&A	20.01.2026	Zoom	Open Lab 4
5	Exam	27.01.2026	HS1 (Linz), S1/S3 (Vienna), S5 (Bregenz)	
9	Retry Exam	24.02.2026	tba	

# Research Example: How Robot Dogs See the Unseeable



Peering Florida Giant Bush Cricket



Peering Locust



Peering Quadruped Robot



Metal Fence and Plants




Camera-Equipped Head-Unit



Conventional Image (RGB)

**LMM Vision Reasoning:**  
The image shows some green plants with thin stems and fine leaves growing in front of a metallic fence or railing. Behind the plants, there is an industrial or commercial-looking structure and a concrete floor. The setting appears to be an industrial facility or a technical building exterior.




SA Integral Image (RGB)

**LMM Vision Reasoning:**  
The image appears to be quite blurry, but some features can still be made out. It seems to show an industrial or abandoned facility, possibly a warehouse or a factory. In the center, there's a large structure with circular openings or pipes — maybe ventilation ducts, tubes, or barrels stacked in a frame. The surroundings include a metal framework.



Conventional Image (NIR)

**LMM Vision Reasoning:**  
The image appears to show a grayscale or infrared view of a scene with dense vegetation in the foreground — mainly leafy plants or shrubs. In the background, there seems to be a structure or series of buildings, possibly with a road or paved area in between.



SA Integral Image (NIR)

**LMM Vision Reasoning:**  
The image appears to be a blurry grayscale or low-resolution photo showing an outdoor scene. It looks like there is some vegetation or bushes in the background, and a person or human-like figure is sitting or crouching on the ground near the center of the image. The surroundings seem natural — possibly a garden, park, or wooded area.

# What does $\star$ do?

$$2 \star 3 = 36$$

$$7 \star 1 = 49$$

$$5 \star 2 = 100$$

$$2 \star 2 = 16$$

# What does $\star$ do?



Training

$$2 \star 3 = 36$$

$$7 \star 1 = 49$$

$$5 \star 2 = 100$$


$$2 \star 2 = 16$$



$$x \star y \rightarrow (xy)^2$$

Testing

$$3 \star 5 \rightarrow$$




$$\rightarrow 225$$


# Machine Learning

## ML Classes:


- Unsupervised
- Supervised
- Semi-supervised
- Self-supervised



# Unsupervised Learning




# Supervised Learning





Use labeled Data for a Prediction Task.

# Semi-Supervised Learning



# Self-Supervised Learning


Dataset



No external Labels. Instead, uses Properties of the Data itself as supervisory Signals. Example (learn how to solve a puzzle): take unlabeled images, rotate them, and train a model to guess the rotation. It needs to understand how the objects look upright → helps to learn how to classify.

# Dimensionality Reduction


Given Data Points in  $d$  Dimensions, convert them to Data Points in  $k < d$  Dimensions with minimal loss of Information.




Original  
Data in 3D

Same data  
in 2D

# Dimensionality Reduction: Example




# What is the best Direction?



What is the “best” direction to represent this data:  $v_1$ ,  $v_2$ , or  $v_3$  ?

Why?

# What is the best Direction?



$v_3$  explains most of the **Variance** in the Data.

Points are most “spread out” in this Direction.

# Principle Component Analysis (PCA)

The Variance along  $v$  of all the Point Projections is:

$$\begin{aligned} \text{var}(v) &= \sum_x \|(x - \bar{x})^T \cdot v\|^2 \\ &= \sum_x v^T (x - \bar{x})(x - \bar{x})^T v \\ &= v^T \left[ \sum_x (x - \bar{x})(x - \bar{x})^T \right] v \\ &= v^T C v , \text{ where } C = X^T X \text{ is the covariance matrix of } X \end{aligned}$$

$v$  = Eigenvectors of  $C$

$\text{var}(v)$  = Eigenvalues of  $v$

PCA is used to compute  
 $v, \text{var}(v)$

We are interested in  $v$   
with the largest  $\text{var}(v)$ !

This works also in higher  
dimensions.

# Example: Eigenfaces




Test Data

Does a similar Image  
of that Person exist?



Training Data

# Example: Eigenfaces



Here, an Image represents a Point in a High Dimensional space


- An  $N \times M$  (binary) Image is a Point in  $R^{NM}$
- We can define Vectors in this Space

# Face Space

- When viewed as Vectors of Pixel Values, Face Images are extremely high-dimensional
  - $100 \times 100$  Image = 10,000 Dimensional
- But very few 10,000-dimensional Vectors are valid Faces
- Let's find a lower-dimensional Representation using PCA!



# Dimensionality Reduction



The set of Faces is a “Subspace” of the Set of Images


- Suppose it is K-dimensional
- We can find the best Subspace using PCA
- This is like fitting a “Hyperplane” to the Set of Faces spanned by vectors  $\mathbf{v}_1, \mathbf{v}_2, \dots, \mathbf{v}_k$

$$\text{Any Face: } \mathbf{x} \approx \bar{\mathbf{x}} + a_1 \mathbf{v}_1 + a_2 \mathbf{v}_2 + \dots + a_k \mathbf{v}_k$$

# Eigenfaces

PCA extracts a set of vectors  $v_1, v_2, v_3, \dots$

Each one of these Vectors is a Direction  
in Face space. What do these look like?



# Projecting onto Eigenfaces


The Eigenfaces  $\mathbf{v}_1, \dots, \mathbf{v}_K$  span the Space of Faces. A Face is converted to Eigenface coordinates by

$$\mathbf{x} \rightarrow ((\underbrace{(\mathbf{x} - \bar{\mathbf{x}}) \cdot \mathbf{v}_1}_{a_1}, \underbrace{(\mathbf{x} - \bar{\mathbf{x}}) \cdot \mathbf{v}_2}_{a_2}, \dots, \underbrace{(\mathbf{x} - \bar{\mathbf{x}}) \cdot \mathbf{v}_K}_{a_K})$$

$$\mathbf{x} \approx \bar{\mathbf{x}} + a_1 \mathbf{v}_1 + a_2 \mathbf{v}_2 + \dots + a_K \mathbf{v}_K$$



# Choosing the Dimension K



How many Eigenfaces to use? Look at the Decay of the Eigenvalues

- The Eigenvalue tells you the amount of Variance “in the Direction” of that Eigenface
- Ignore Eigenfaces with low Variance


# Face Recognition with Eigenfaces

- Project Input Face to Eigenfaces (i.e., represent it by a K-dimensional vector:  $\mathbf{a}_1, \mathbf{a}_2, \dots, \mathbf{a}_K$ )
- Find the Nearest-Neighbor in Training Image Set (also projected to Eigenfaces)
- Pro: Speed, Size, Fairly Invariant to Resolution
- Cons: Fails quickly if Variation between Training Images and Test Image becomes too large

# Dimensionality Reduction Methods

- PCA (Principal Component Analysis):
  - Find Projection that maximize the Variance
- ICA (Independent Component Analysis):
  - Very similar to PCA except that it assumes non-Gaussian features
- Multidimensional Scaling:
  - Find Projection that best preserves Inter-Point Distances
- LDA(Linear Discriminant Analysis):
  - Maximizing the Component Axes for Class-Separation
- ...
- ...

# Supervised Learning – Example: Image Classification



Use labeled Data for a Prediction Task.

# Learning from Labeled Training Data


Training data

$$\{x^{(1)}, y^{(1)}\}$$

$$\{x^{(2)}, y^{(2)}\} \rightarrow$$

$$\{x^{(3)}, y^{(3)}\}$$





...



$$\rightarrow f : X \rightarrow Y$$

# Linear Regression

Training data




$$f_{\theta}(x) = \theta_1 x + \theta_0$$

Hypothesis space

The relationship between X and Y is roughly linear:  $y \approx \theta_1 x + \theta_0$

# Linear Regression

Training data



Search for the Parameters,  $\theta = \{\theta_0, \theta_1\}$  ,  
that best fit the Data.

$$f_\theta(x) = \theta_1 x + \theta_0$$


**Best Fit in what sense?**

The Least-Squares **Objective** (aka **Loss**)  
says the best Fit is the Function that  
minimizes the Squared Error between  
Predictions and Target Values:

$$\mathcal{L}(\hat{y}, y) = (\hat{y} - y)^2 \quad \hat{y} \equiv f_\theta(x)$$

# Linear Regression



Data  
 $\{x^{(i)}, y^{(i)}\}_{i=1}^N \rightarrow$





# Linear Regression

Training

Data





Learner



We could assume more complex Functions (like Polynomial Regression) – but still very limited in terms of **generalization**...

Testing

Input






Output

# Deep Learning


Training

Data




Testing

Input →




→ Output

# Deep Learning



$$\theta^* = \arg \min_{\theta} \sum_{i=1}^N \mathcal{L}(f_{\theta}(\mathbf{x}^{(i)}), \mathbf{y}^{(i)})$$

# Deep Learning



**Training:** Used to train Model Parameters


**Validation:** Used to tune Hyperparameters (e.g.,  $\lambda$  in regularized regression)

**Testing:** Used to report final Accuracy (shouldn't be touched before!)



# Classification with Convolution





What's the Object Class of the Center Pixel?





“Bird”



“Bird”




“Sky”






“Sky”


# Classification with Convolution



# Classification with Convolution



**W** computes a weighted Sum of all Pixels in the Patch




**W** is a Convolutional Kernel applied to the Full Image!

It gives highest Response if the Kernel matches the Image Content

It is Shift Invariant. But how about Scale, Rotation, other Poses, etc.?


# Convolutional Neural Networks (CNNs)

Input to CNN layer  $l$


$$x_l \in R^{H \times W \times C_l}$$


# Convolutional Layer

$32 \times 32 \times 3$  image




$5 \times 5 \times 3$  filter






Let's apply a different filter  
with different weights!

Activation maps




# Activation (Feature) Maps



- Each Layer can be thought of as a set of C **Feature Maps** aka **Channels**
- Each Feature Map is an NxM Image
- ReLU (Rectified Linear Unit,  $f(x)=\max(0,x)$ ) Activation Function is used to introduce Nonlinearity in a Neural Network, helping mitigate the Vanishing Gradient Problem


# Multiple Convolutional Layers




# Pooling

**Max pooling**  
(usually 2x2 or 3x3)


$$y_j = \max_{j \in \mathcal{N}(j)} h_j$$



provides large Response  
regardless of exact Position of  
Edge




# Pooling and Down-Sampling




**Stride:** shift of kernel (for normal Convolution it is 1)



# Computations in a CNN



# Computations in a CNN



# Receptive Field



**Receptive Field** is the Region of the Image that a particular Neuron in a Convolutional Layer is “looking at” or taking into account when making its Predictions or Feature Extractions (depends on Kernel Size, Stride, Pooling).

# Course Overview

CW	Topic	Date	Place	Lab
41	Introduction and Course Overview	07.10.2025	Zoom	Lab 1
42	Capturing Digital Images	14.10.2025	Zoom	Lab 2
43	Digital Image Processing	21.10.2025	Zoom	Assignment 1
44	Machine Learning	28.10.2025	Zoom	
→ 45	Feature Extraction	04.11.2025	Zoom	Open Lab 1
46	Segmentation	11.11.2025	Zoom	Assignment 2
47	Optical Flow	18.11.2025	Zoom	Open Lab 2
48	Object Detection	25.11.2025	Zoom	Assignment 3
49	Multi-View Geometry	02.12.2025	Zoom	Open Lab 3
50	3D Vision	09.12.2025	Zoom	Assignment 4
3	Trends in Computer Vision	13.01.2026	Zoom	
4	Q&A	20.01.2026	Zoom	Open Lab 4
5	Exam	27.01.2026	HS1 (Linz), S1/S3 (Vienna), S5 (Bregenz)	
9	Retry Exam	24.02.2026	tba	

# Next Week: Feature Extraction

## Model-Based vs. Learning-Based Feature Extraction

- Fixed engineered features (or kernels) + trainable classifier




VS.

- End-to-end learning / feature learning / deep learning



**JKU** JOHANNES KEPLER  
UNIVERSITY LINZ

## From Low-Level to High-Level Features



**JKU** JOHANNES KEPLER  
UNIVERSITY LINZ

## Example: Scale-Invariant Feature Transform (SIFT)




(3) Filter out Features in low-contrast Regions  
(Noise)

**JKU** JOHANNES KEPLER  
UNIVERSITY LINZ

(4) Determine Feature (i.e., Gradient)  
Orientations and sort them into Histogram  
(largest Bin = main Orientation)

## Development of Architectures



**JKU** JOHANNES KEPLER  
UNIVERSITY LINZ

# Thank You

