

Computer Vision



Lecture 6: Segmentation
Oliver Bimber

Last Week: Feature Extraction

Model-Based vs. Learning-Based Feature Extraction

- Fixed engineered features (or kernels) + trainable classifier



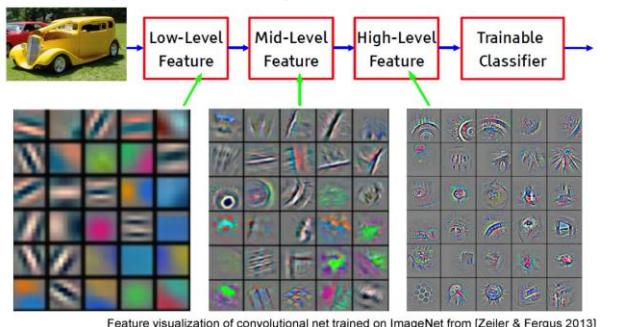
vs.

- End-to-end learning / feature learning / deep learning



JKU JOHANNES KEPLER
UNIVERSITY LINZ

From Low-Level to High-Level Features



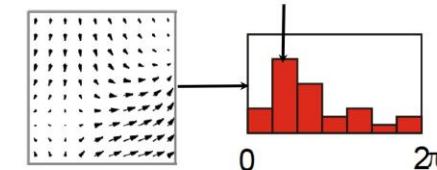
JKU JOHANNES KEPLER
UNIVERSITY LINZ

Example: Scale-Invariant Feature Transform (SIFT)



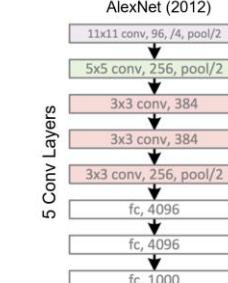
(3) Filter out Features in low-contrast Regions (Noise)

JKU JOHANNES KEPLER
UNIVERSITY LINZ

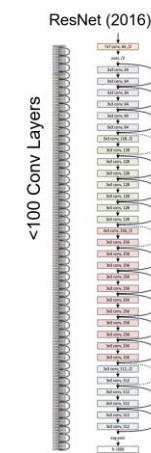
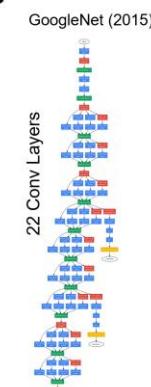
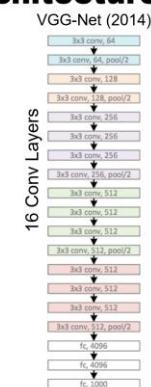


(4) Determine Feature (i.e., Gradient) Orientations and sort them into Histogram (largest Bin = main Orientation)

Development of Architectures



...going really deep..



Course Overview

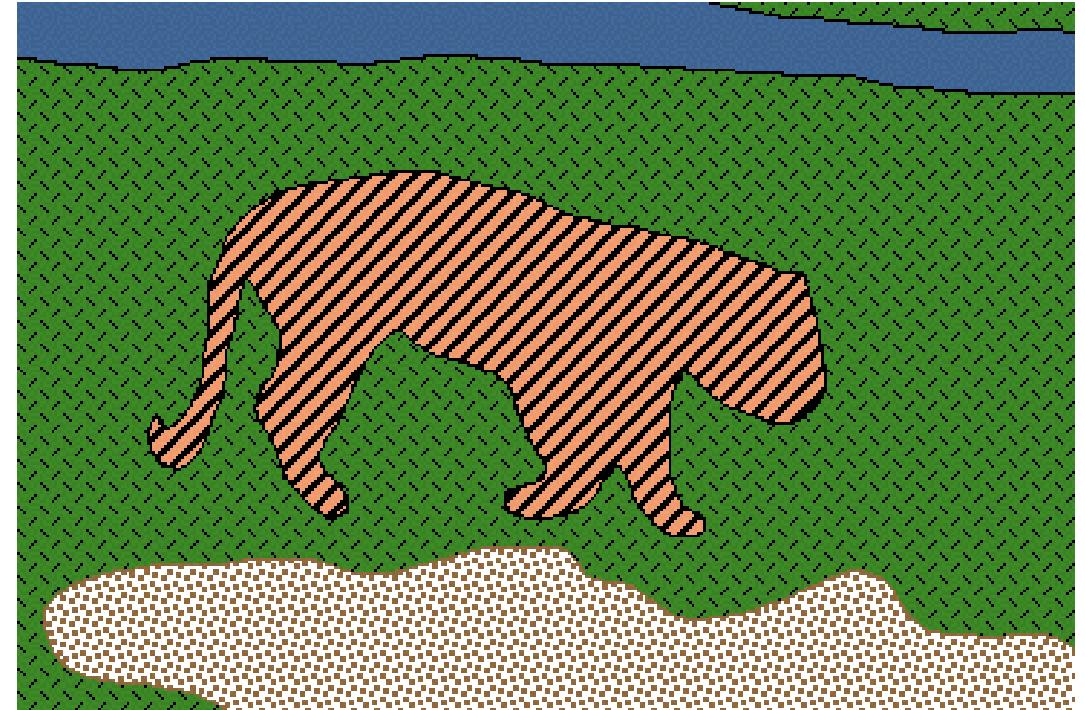
CW	Topic	Date	Place	Lab
41	Introduction and Course Overview	07.10.2025	Zoom	Lab 1
42	Capturing Digital Images	14.10.2025	Zoom	Lab 2
43	Digital Image Processing	21.10.2025	Zoom	Assignment 1
44	Machine Learning	28.10.2025	Zoom	
45	Feature Extraction	04.11.2025	Zoom	Open Lab 1
→ 46	Segmentation	11.11.2025	Zoom	Assignment 2
47	Optical Flow	18.11.2025	Zoom	Open Lab 2
48	Object Detection	25.11.2025	Zoom	Assignment 3
49	Multi-View Geometry	02.12.2025	Zoom	Open Lab 3
50	3D Vision	09.12.2025	Zoom	Assignment 4
3	Trends in Computer Vision	13.01.2026	Zoom	
4	Q&A	20.01.2026	Zoom	Open Lab 4
5	Exam	27.01.2026	HS1 (Linz), S1/S3 (Vienna), S5 (Bregenz)	
9	Retry Exam	24.02.2026	tba	

Research Examples:

Today 1:45pm, HS1 JKU, or:

<https://jku.zoom.us/j/97166662151?pwd=wolaNes9BclitV6Vui32jBHTIyJ4VA.1>

What is Segmentation?

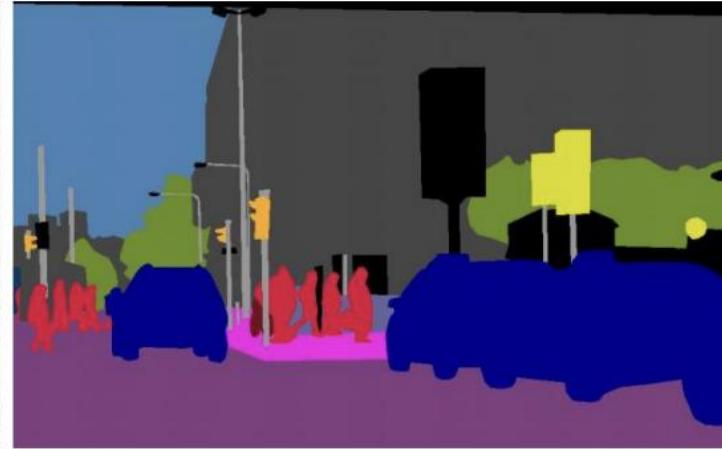


Identify Groups of Pixels that go together

Types of Segmentation



Semantic (by Classes)



Instance (by Objects)



Panoptic (by Classes and Objects)

Classification vs. Segmentation

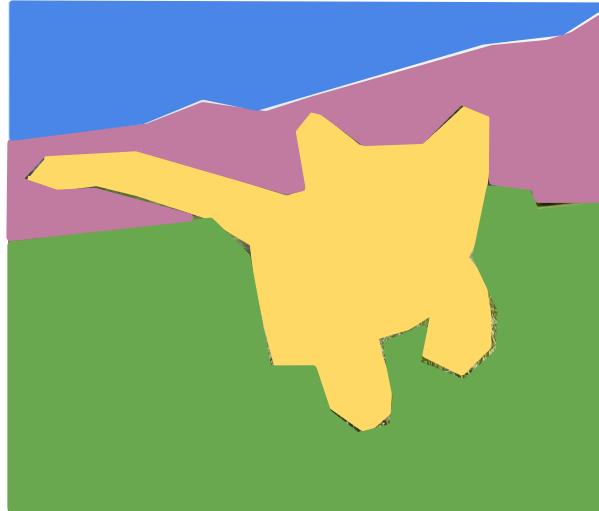
Classification



CAT

No spatial extent

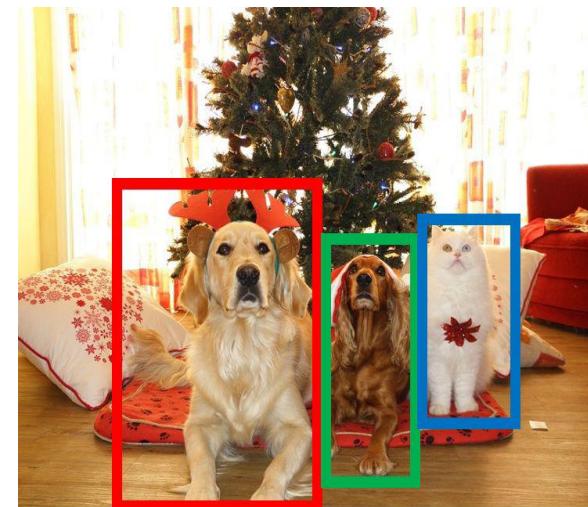
Semantic
Segmentation



GRASS, CAT,
TREE, SKY

No objects, just pixels

Object
Detection



DOG, DOG, CAT

Multiple Object

Instance
Segmentation



DOG, DOG, CAT

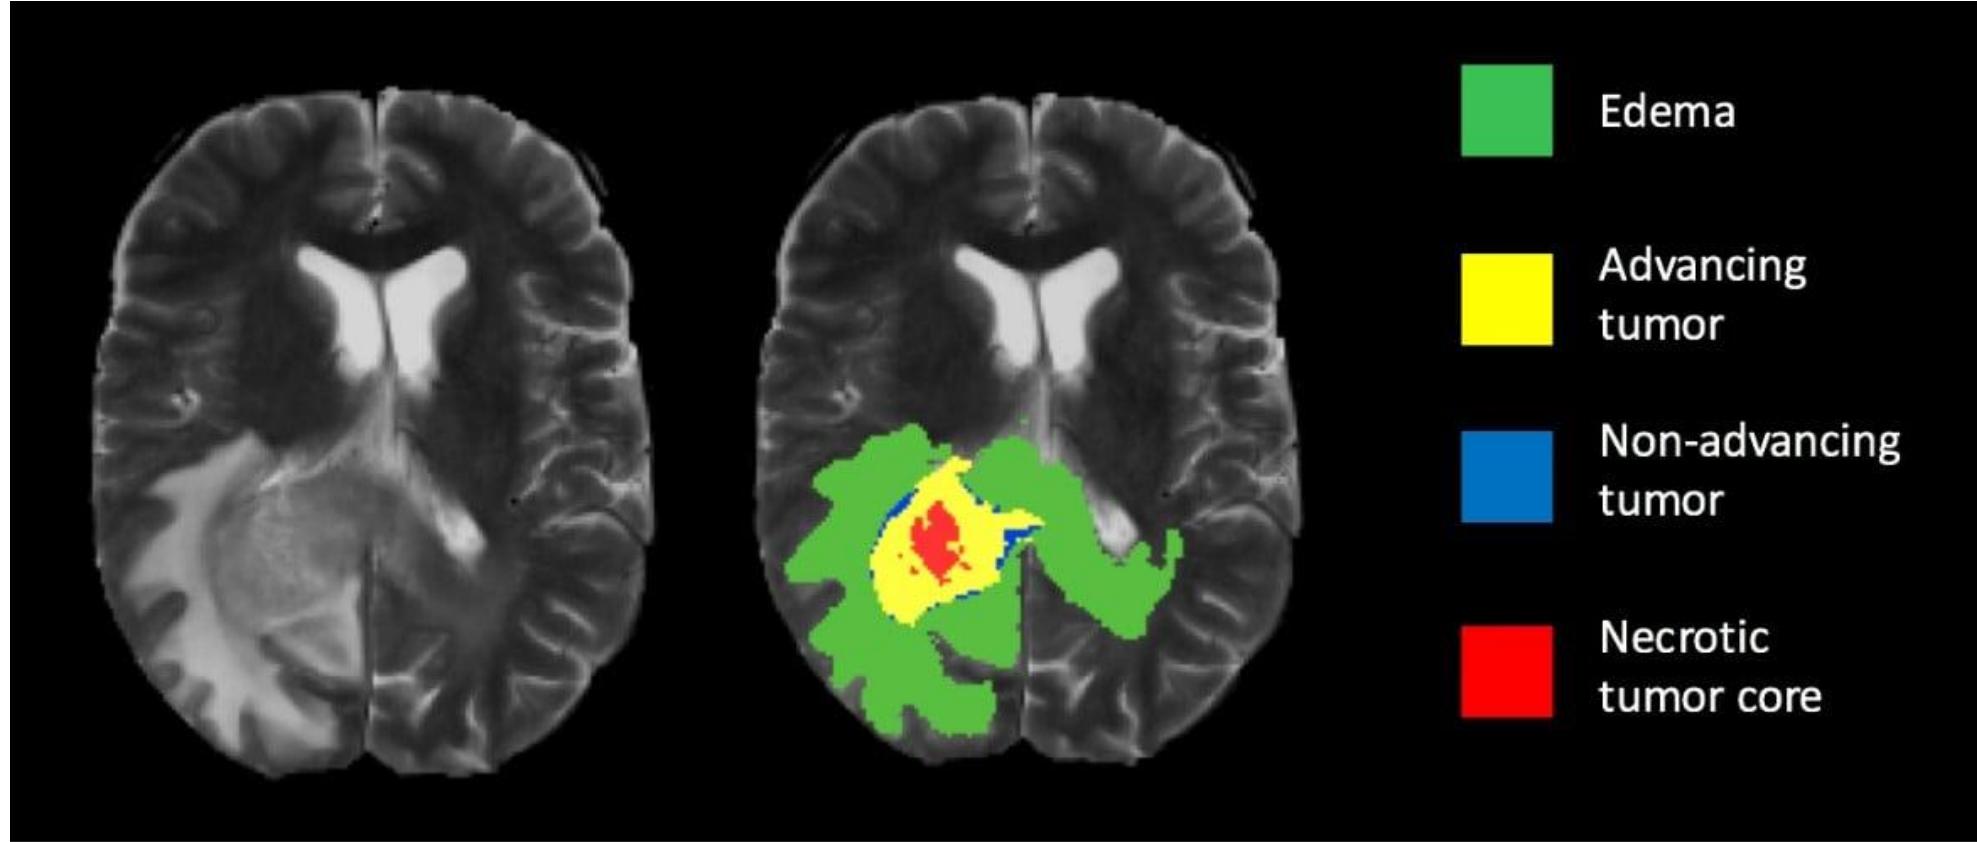
Example: Image Editing



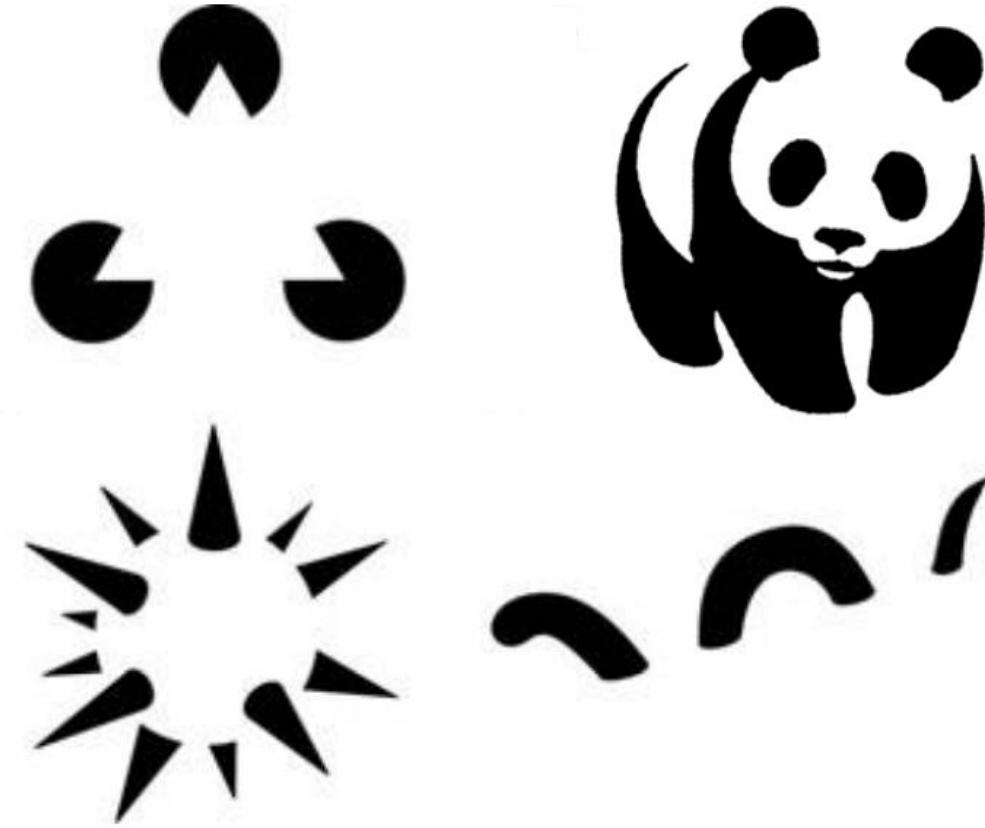
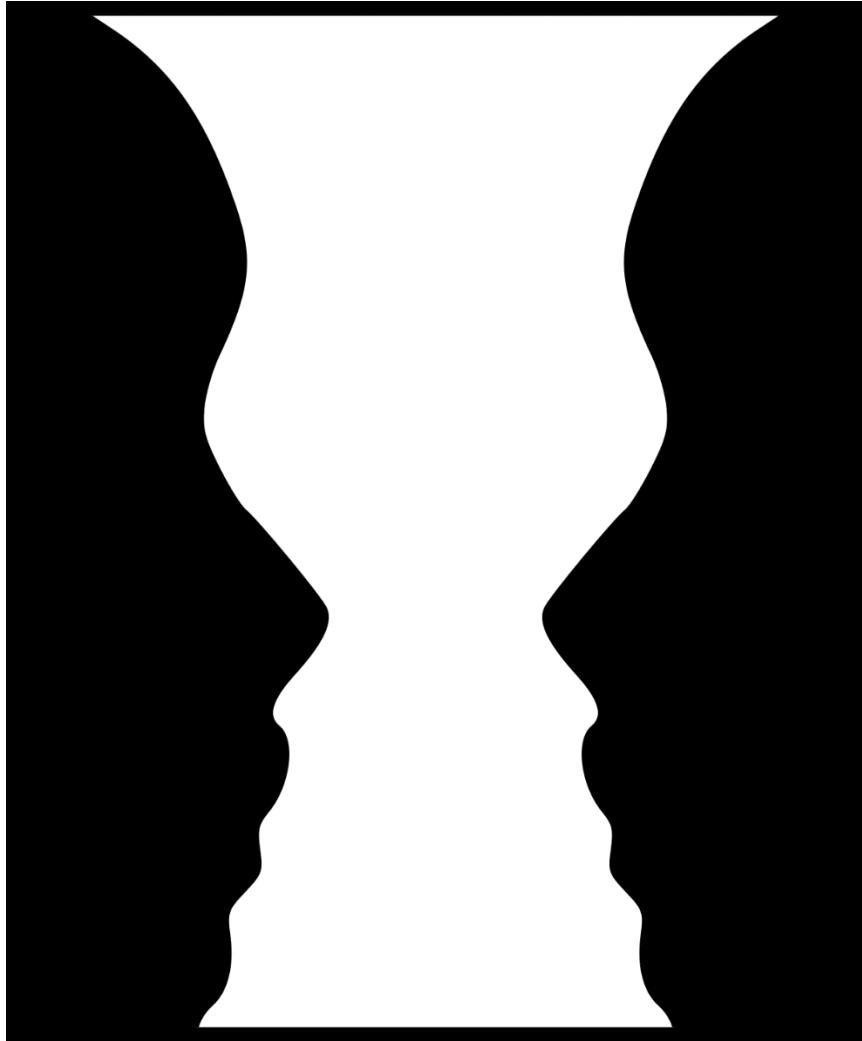
Example: Autonomous Driving



Example: Medical Imaging

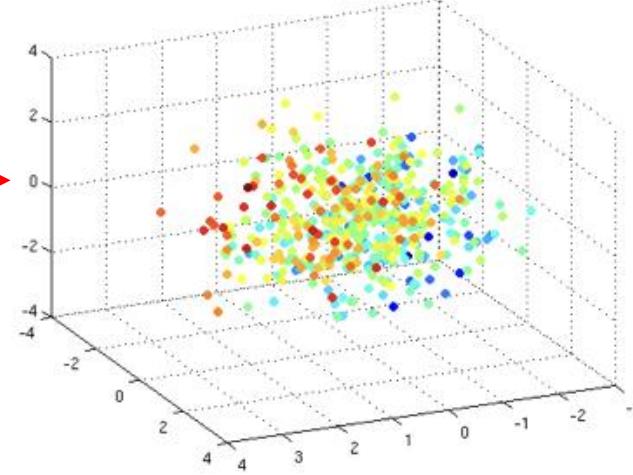


Segmentation done by Humans



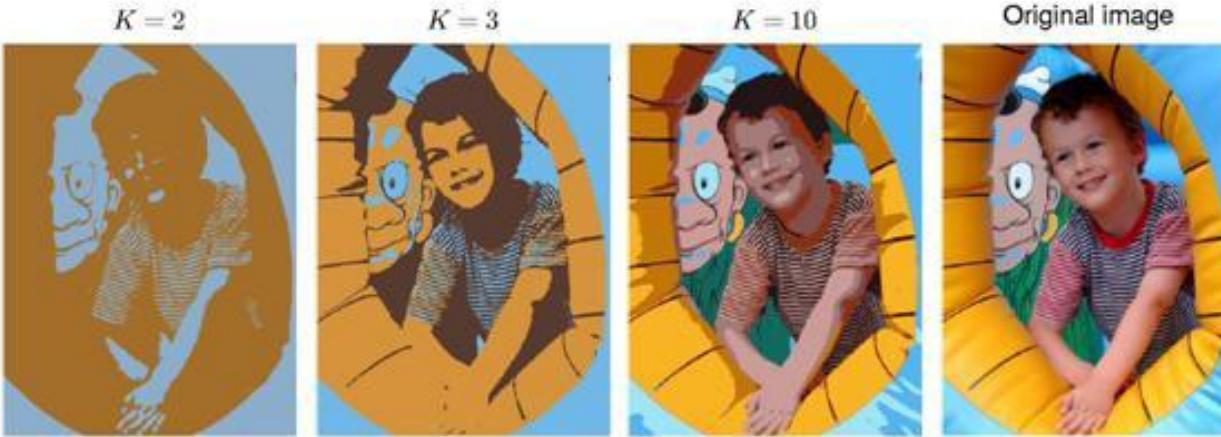
Illusory or subjective contours are perceived

Segmentation by Clustering



- Pixels are Points in a high-dimensional Space, eg.:
 - Color: 3d
 - Color + Location: 5d
- Cluster Pixels into Segments, eg.:
 - K-Means Clustering

Example: K-Means Clustering



$$\Phi_{(\text{cluster}, \text{data})} = \sum_{i \in \text{cluster}} \left\{ \sum_{j \in \text{cluster}(i)} (x_j - c_i)^T (x_j - c_i) \right\}$$

K-means:

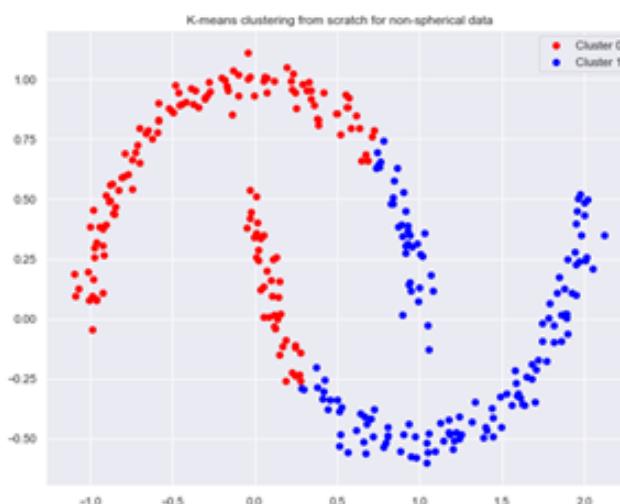
choose k points by random to act as cluster centers

until cluster centers are unchanged

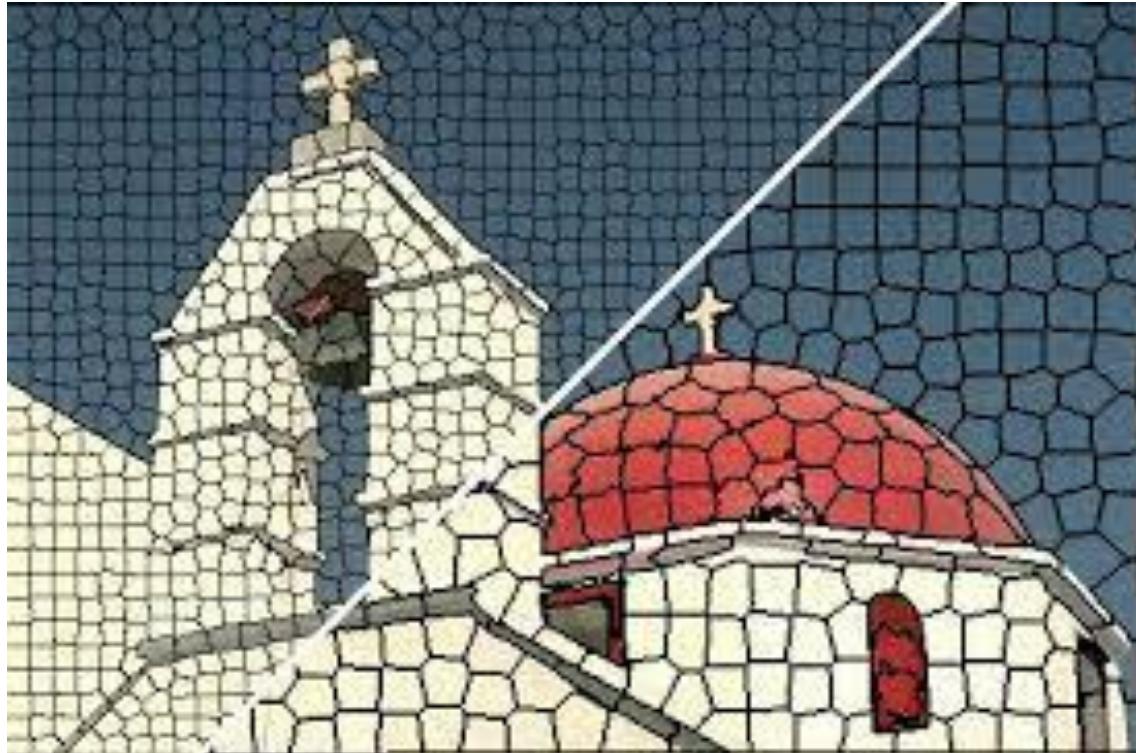
 allocate each point to nearest cluster
 (ensure that each cluster has at least one point)

 replace cluster centers with mean of new cluster points

end



Example: K-Means Clustering



As K increases...

$$\Phi_{(\text{cluster}, \text{data})} = \sum_{i \in \text{cluster}} \left\{ \sum_{j \in \text{cluster}(i)} (x_j - c_i)^T (x_j - c_i) \right\}$$

K-means:

choose k points by random to act as cluster centers

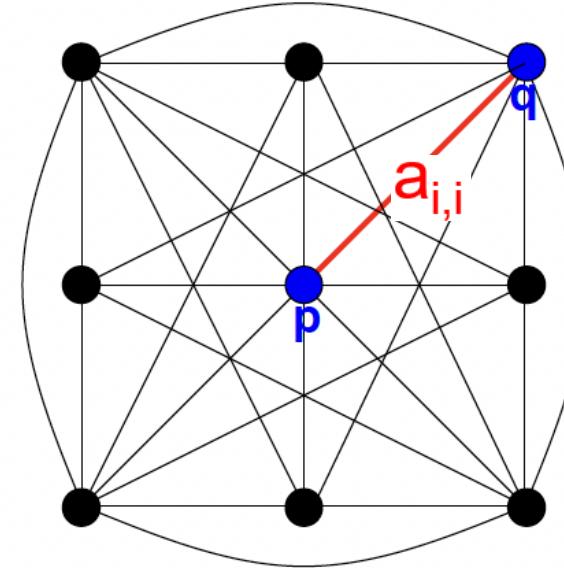
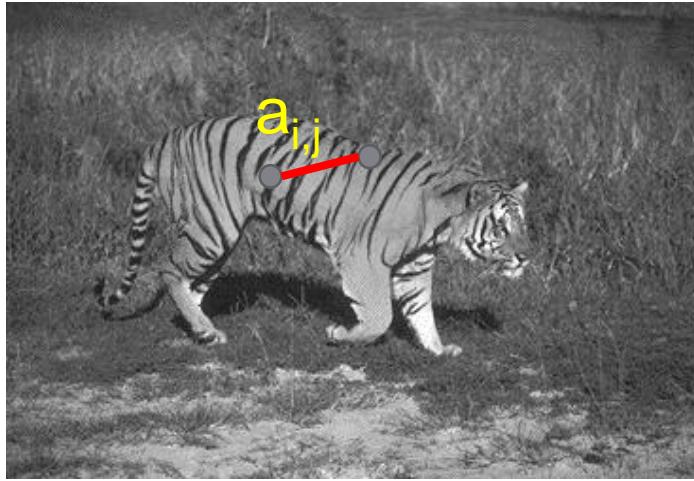
until cluster centers are unchanged

 allocate each point to nearest cluster
 (ensure that each cluster has at least one point)

 replace cluster centers with mean of new cluster points

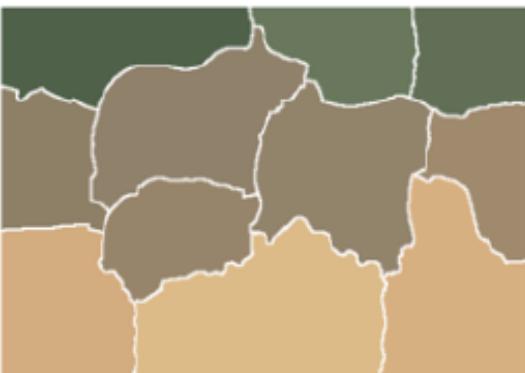
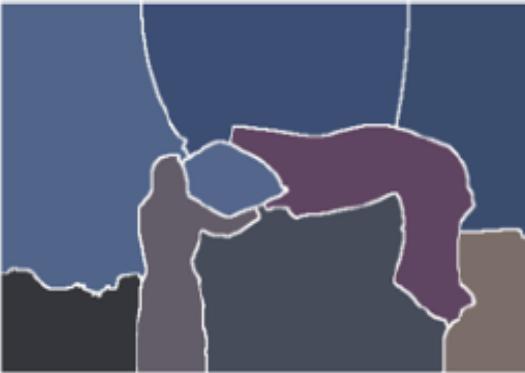
end

Segmentation with Graphs



- Pixels are Nodes in a Graph
 - Edge between Pairs of Pixels (i,j)
 - Affinity Weight $a_{i,i}$ for each Edge measures “Similarity”
- Cluster Pixels into Segments, eg.:
 - Clustering by Graph Eigenvectors, Graph Cut, Grab Cut

Example: Clustering by Graph Eigenvectors



Let's assume your Graph is represented with an Affinity Matrix \mathbf{A} (size: $\mathbf{N} \times \mathbf{N}$ for \mathbf{N} Elements)

A good Cluster is one where Elements that are strongly connected to the Cluster also have large Values connecting one another in the Affinity Matrix

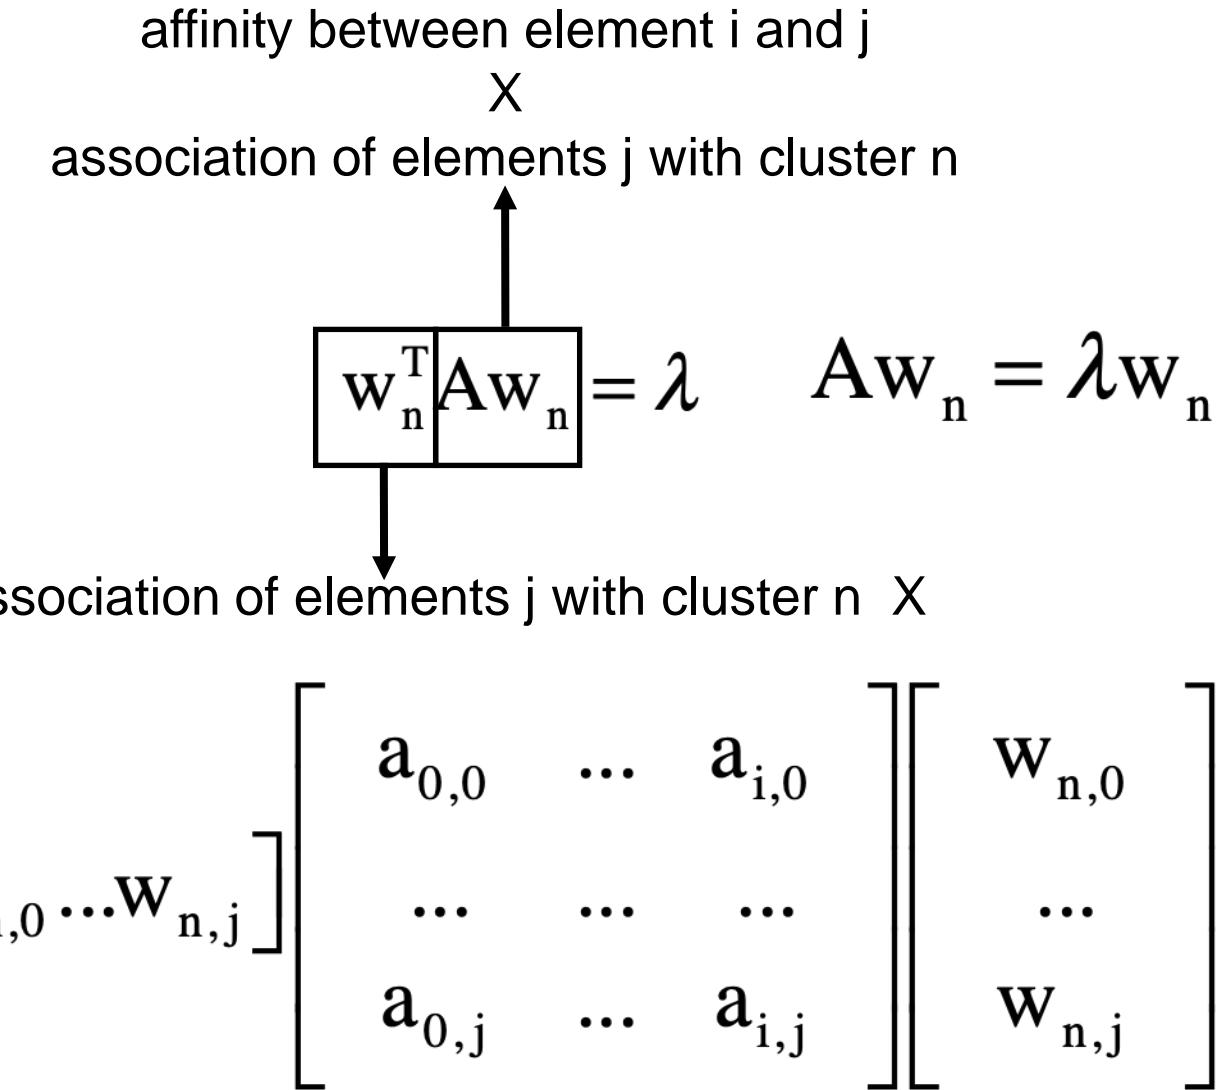
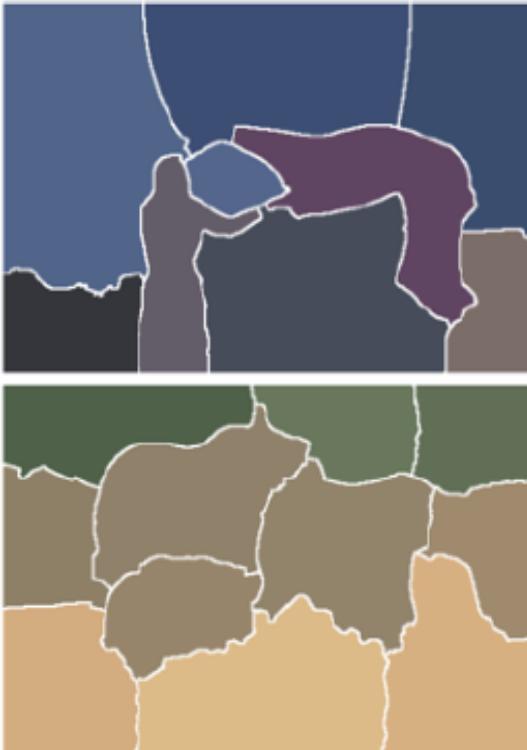
Let's assume you make an Assignment for Cluster n using a Weight Vector \mathbf{w}_n (with N elements, high values indicate strong connectivity to the Cluster)

When is use Assignment a good one?

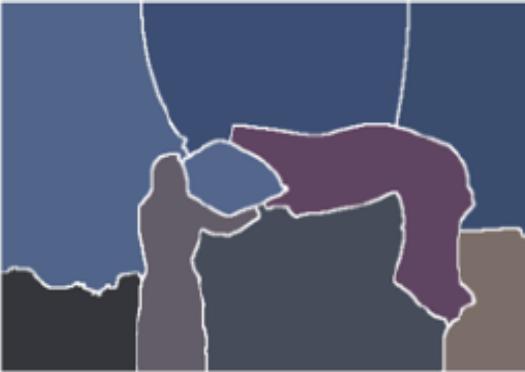
Our objective from above rephrased mathematically:

$$\text{Max}(\mathbf{w}_n^T \mathbf{A} \mathbf{w}_n = \lambda)$$

Example: Clustering by Graph Eigenvectors

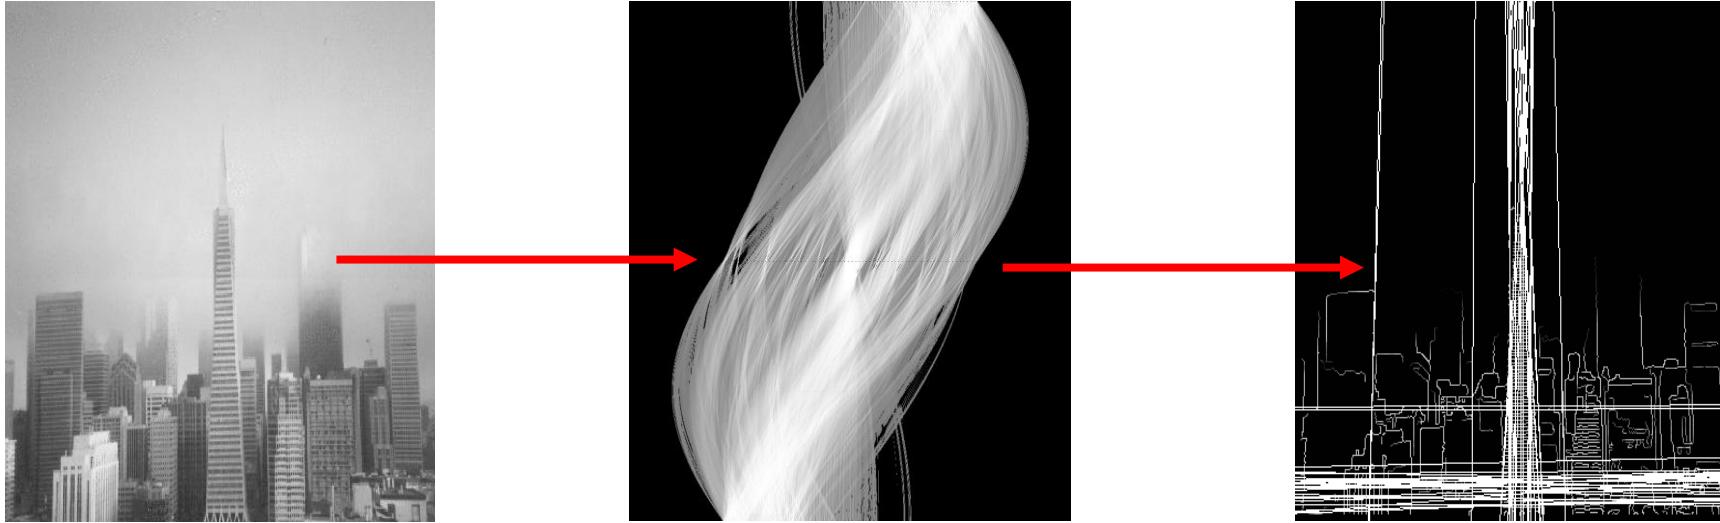


Example: Clustering by Graph Eigenvectors



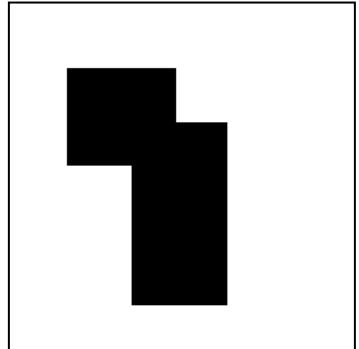
Clustering by Graph Eigenvectors:
construct affinity matrix A
compute eigenvalues and eigenvectors of A
until there are sufficient clusters
 take the eigenvector corresponding to
 next largest eigenvalue
 assign elements to cluster (multiply
 eigenvector with A and threshold)
 zero out all clustered elements in A
end

Segmentation by Fitting

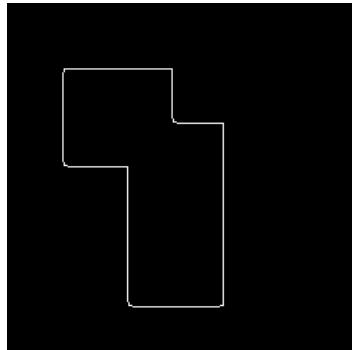


- Knowing the Shapes of the Segments
 - Segments can be parameterized
- Transform Pixels into Parameter Space, eg.:
 - Hough Transform

Example: Hough Transform



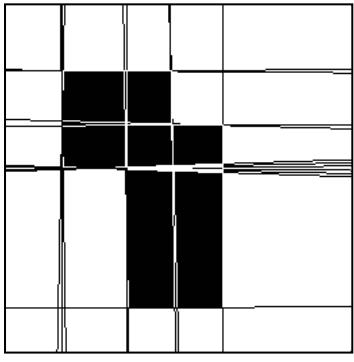
original image



edge detector

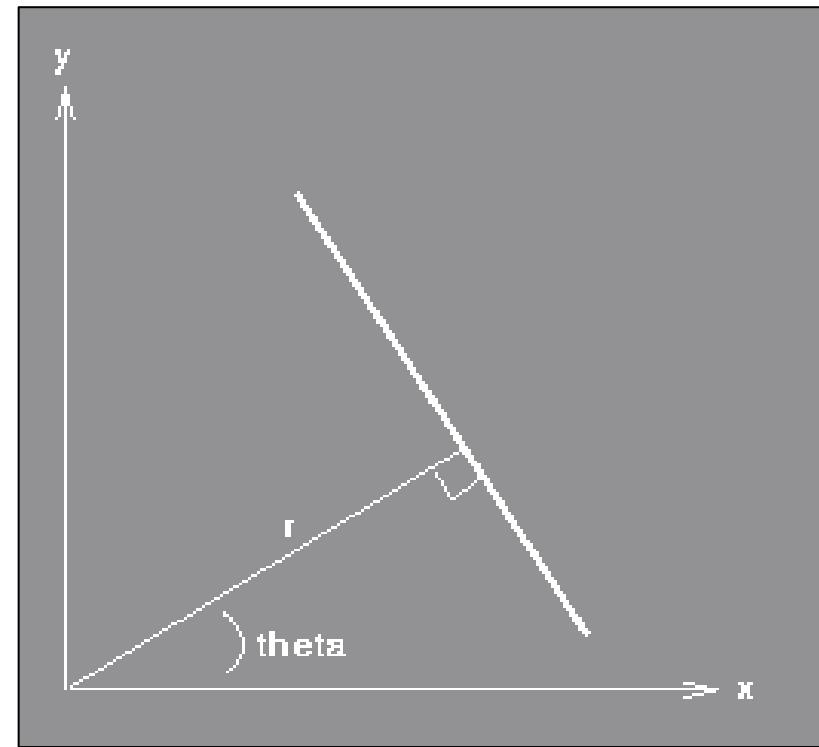


line space



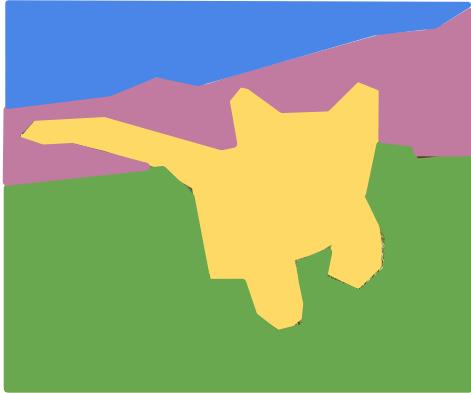
fitted lines

$$x \cos \theta + y \sin \theta + r = 0$$



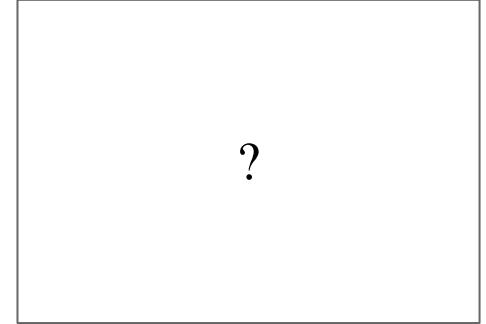
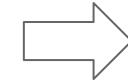
line space parameterization

Segmentation by Learning



**GRASS, CAT,
TREE, SKY, ...**

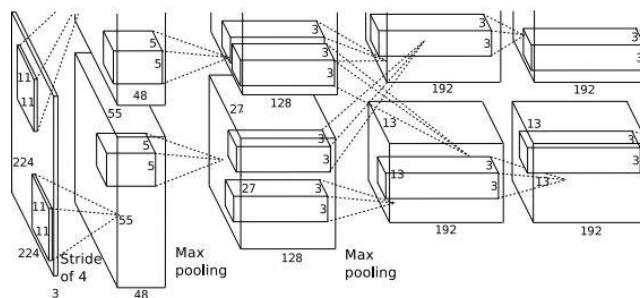
Paired Training Data: for each Training Image,
each Pixel is labeled with a Semantic Category



At Test Time, Classify each Pixel of a new Image

Segmentation using CNNs

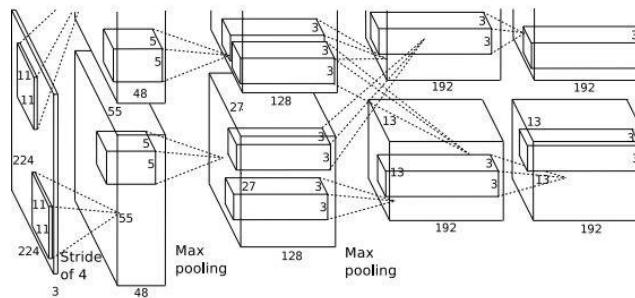
Full image



An intuitive Idea: encode the entire Image with a CNN, and do Semantic Segmentation on top

Segmentation using CNNs

Full image

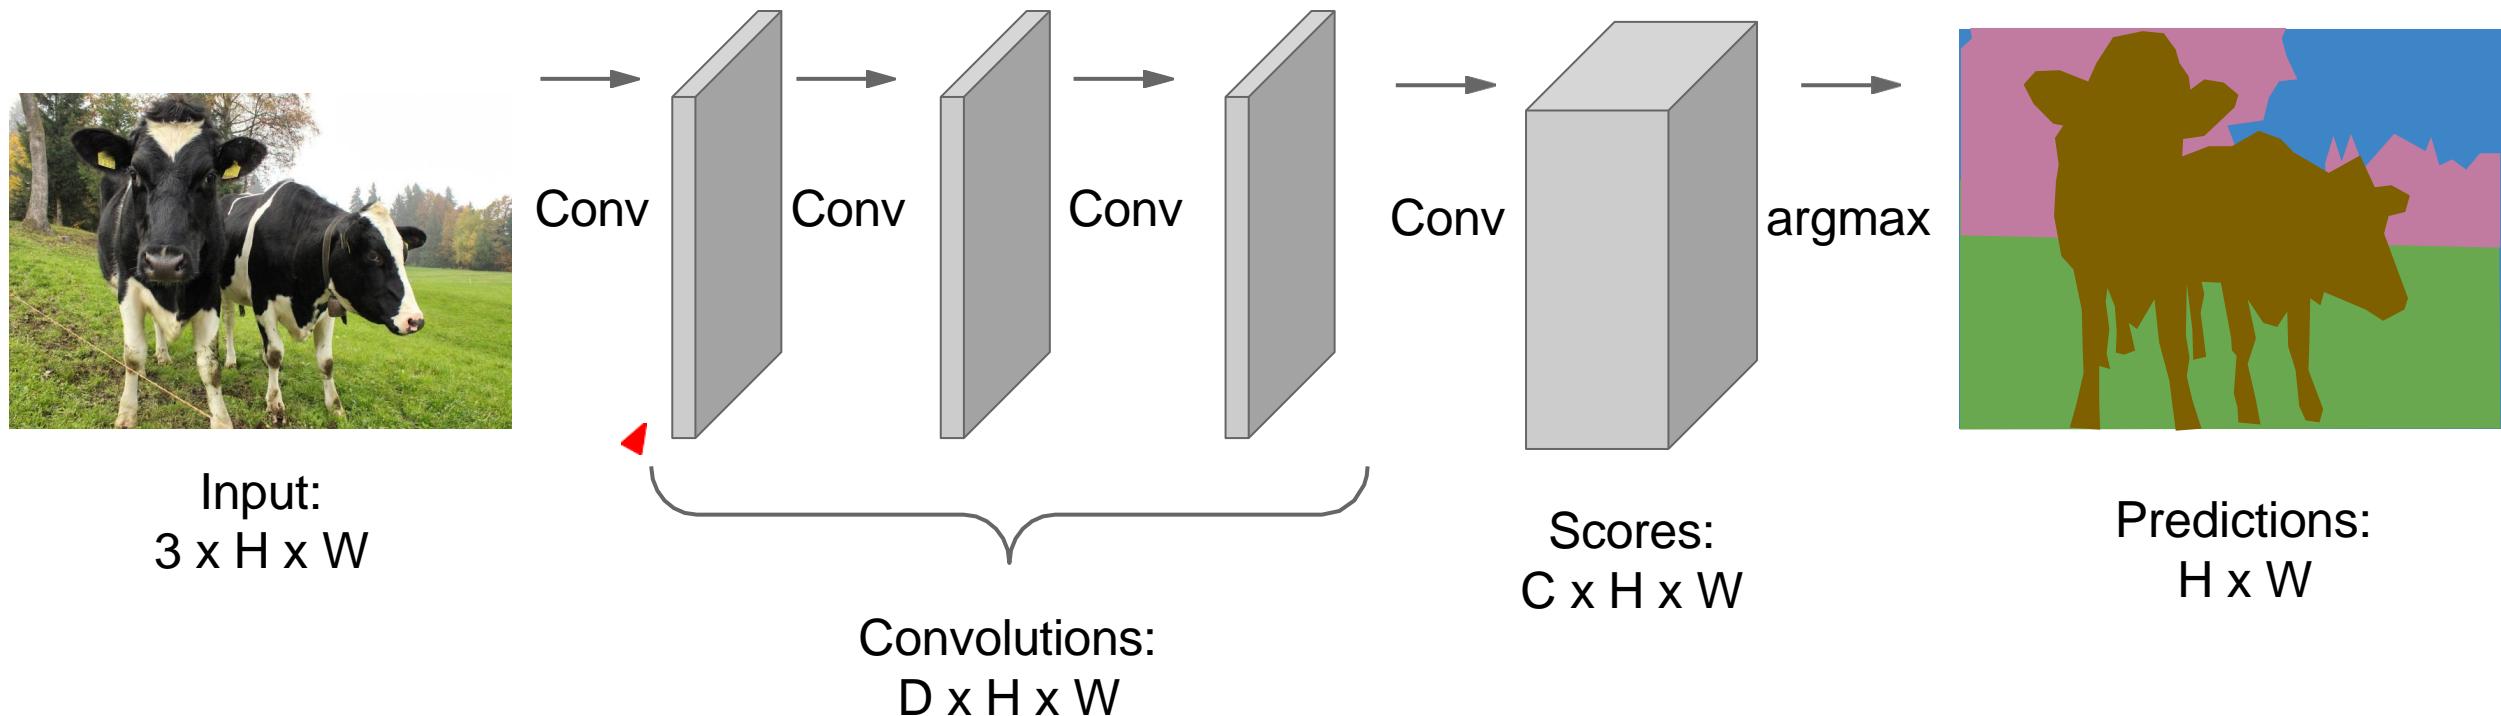


An intuitive Idea: encode the entire Image with a CNN, and do Semantic Segmentation on top

Problem: Classification Architectures often reduce Feature spatial Sizes to go deeper, but Semantic Segmentation requires the Output Size to be the same as Input Size

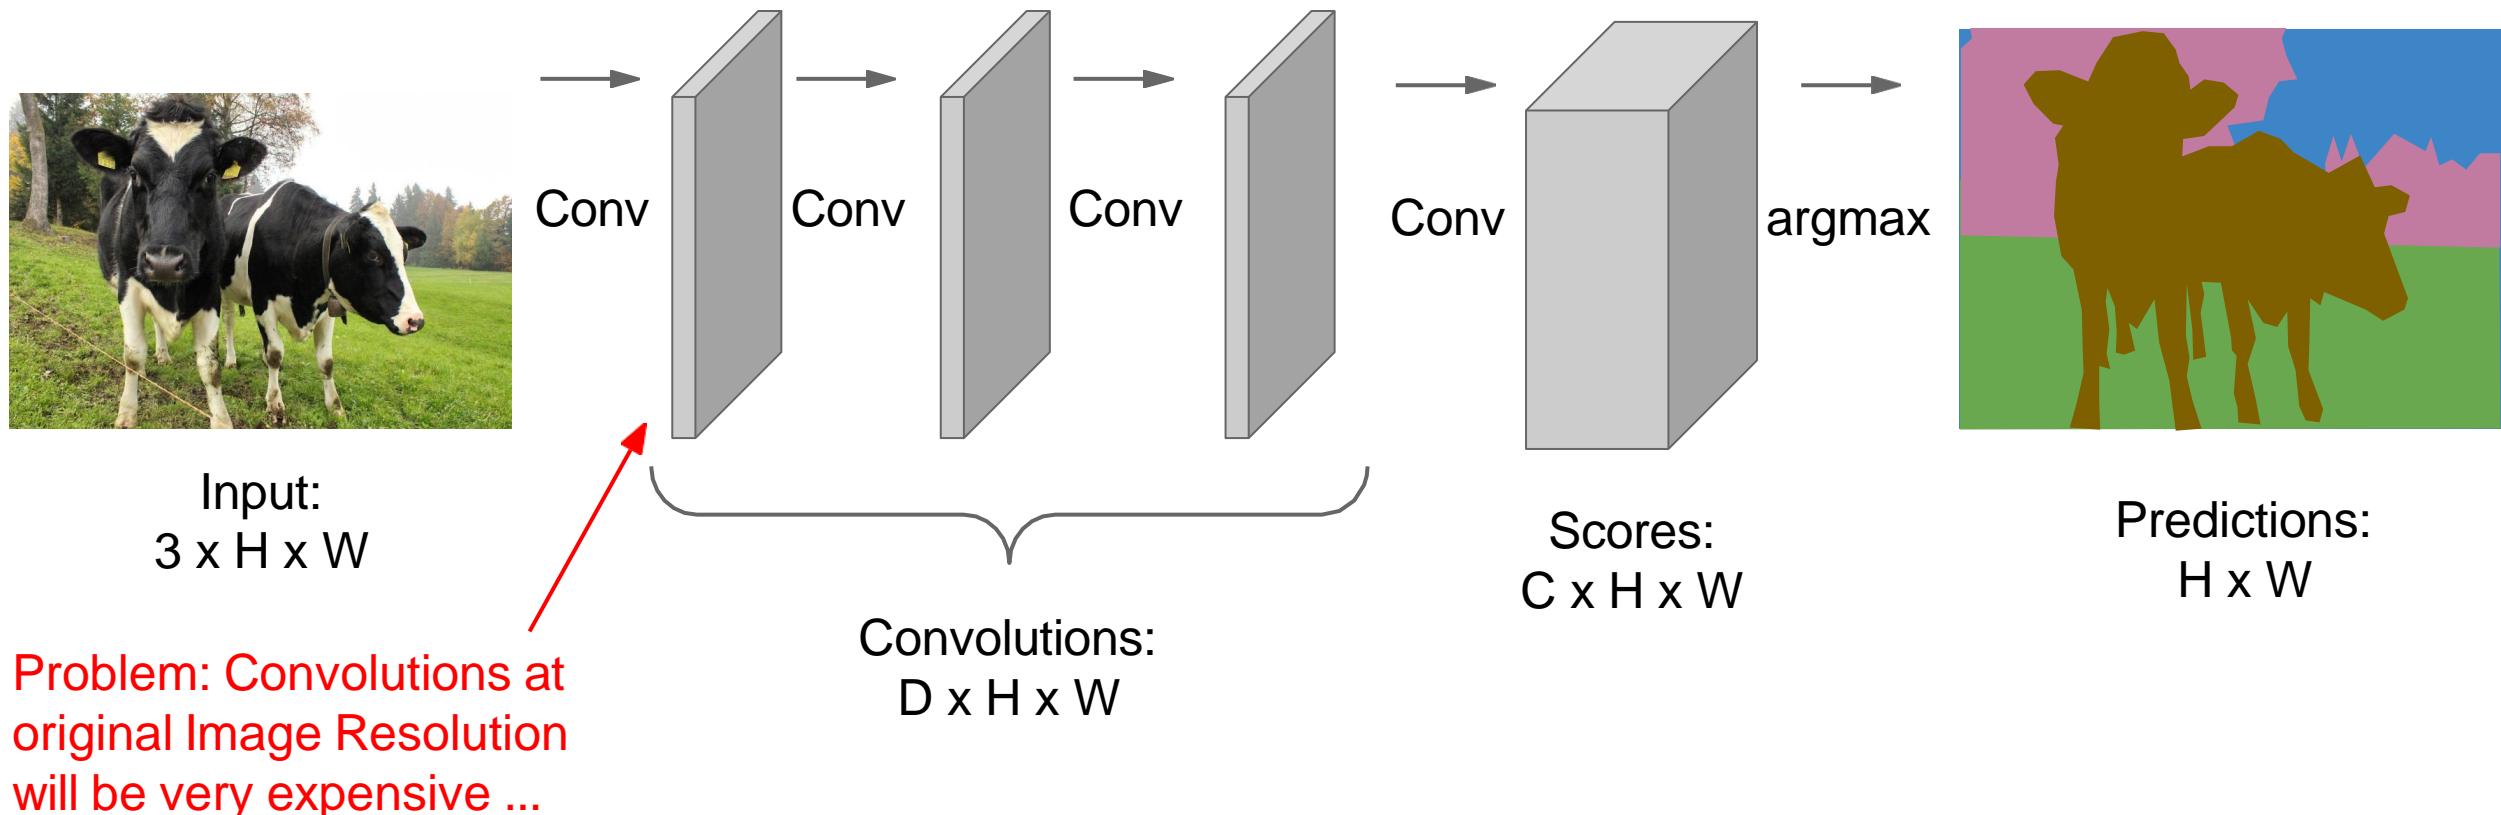
Segmentation using CNNs

Design a Network with only Convolutional Layers without Downsampling Operators to make Predictions for Pixels all at once!



Segmentation using CNNs

Design a Network with only Convolutional Layers without Downsampling Operators to make Predictions for Pixels all at once!



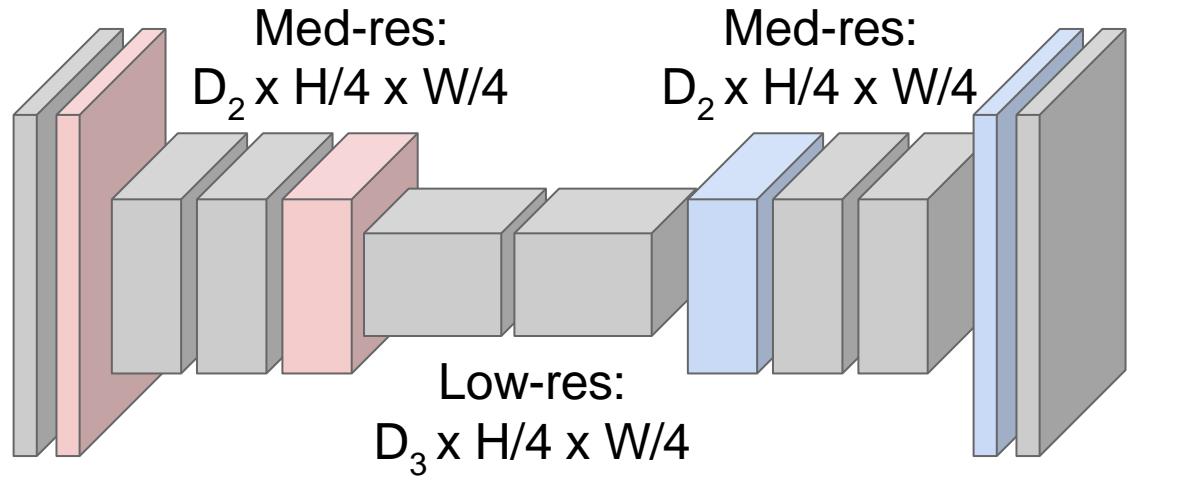
Segmentation using CNNs

Downsampling:
Pooling, Strided
Convolution



Input:
 $3 \times H \times W$

High-res:
 $D_1 \times H/2 \times W/2$

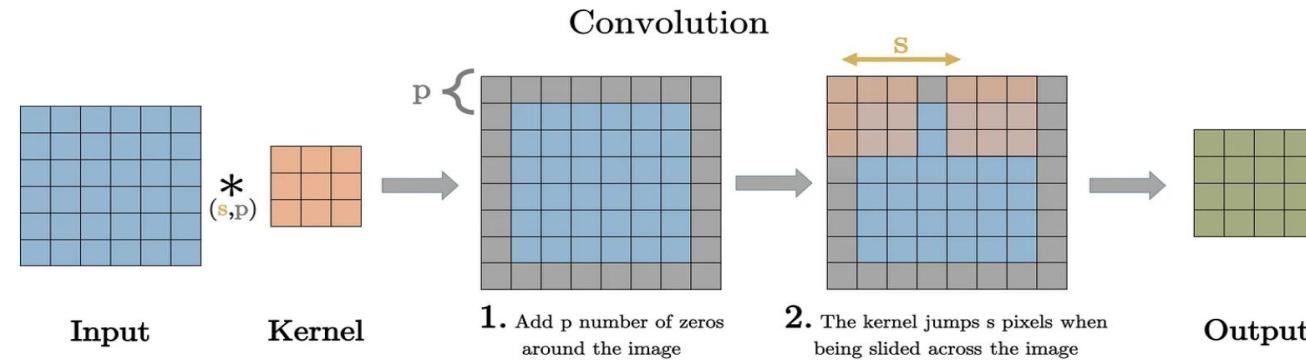
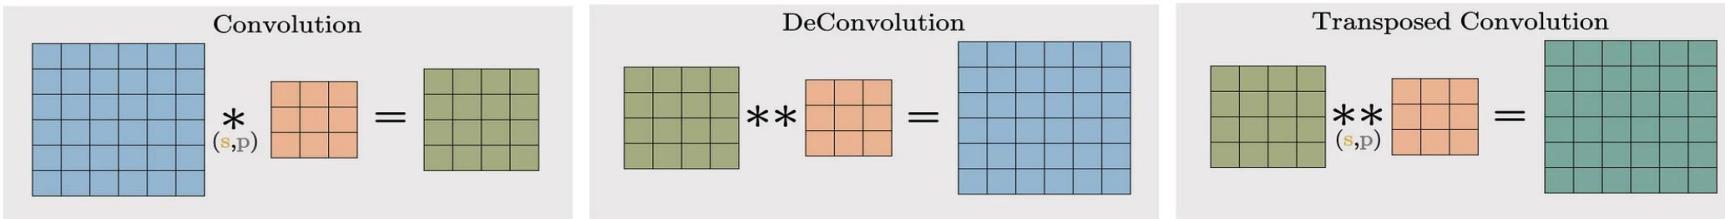


Design Network as a Bunch of Convolutional Layers, with
Downsampling and **Upsampling** inside the Network!

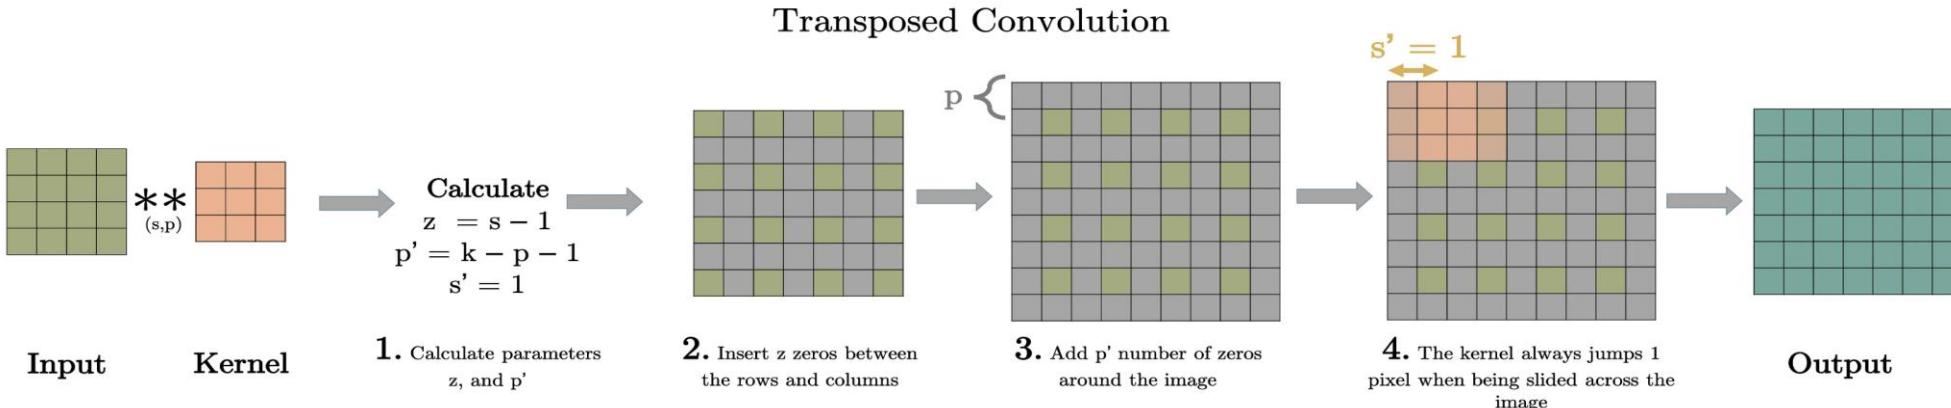
Upsampling:
Strided Transposed
Convolution (or Unpooling)



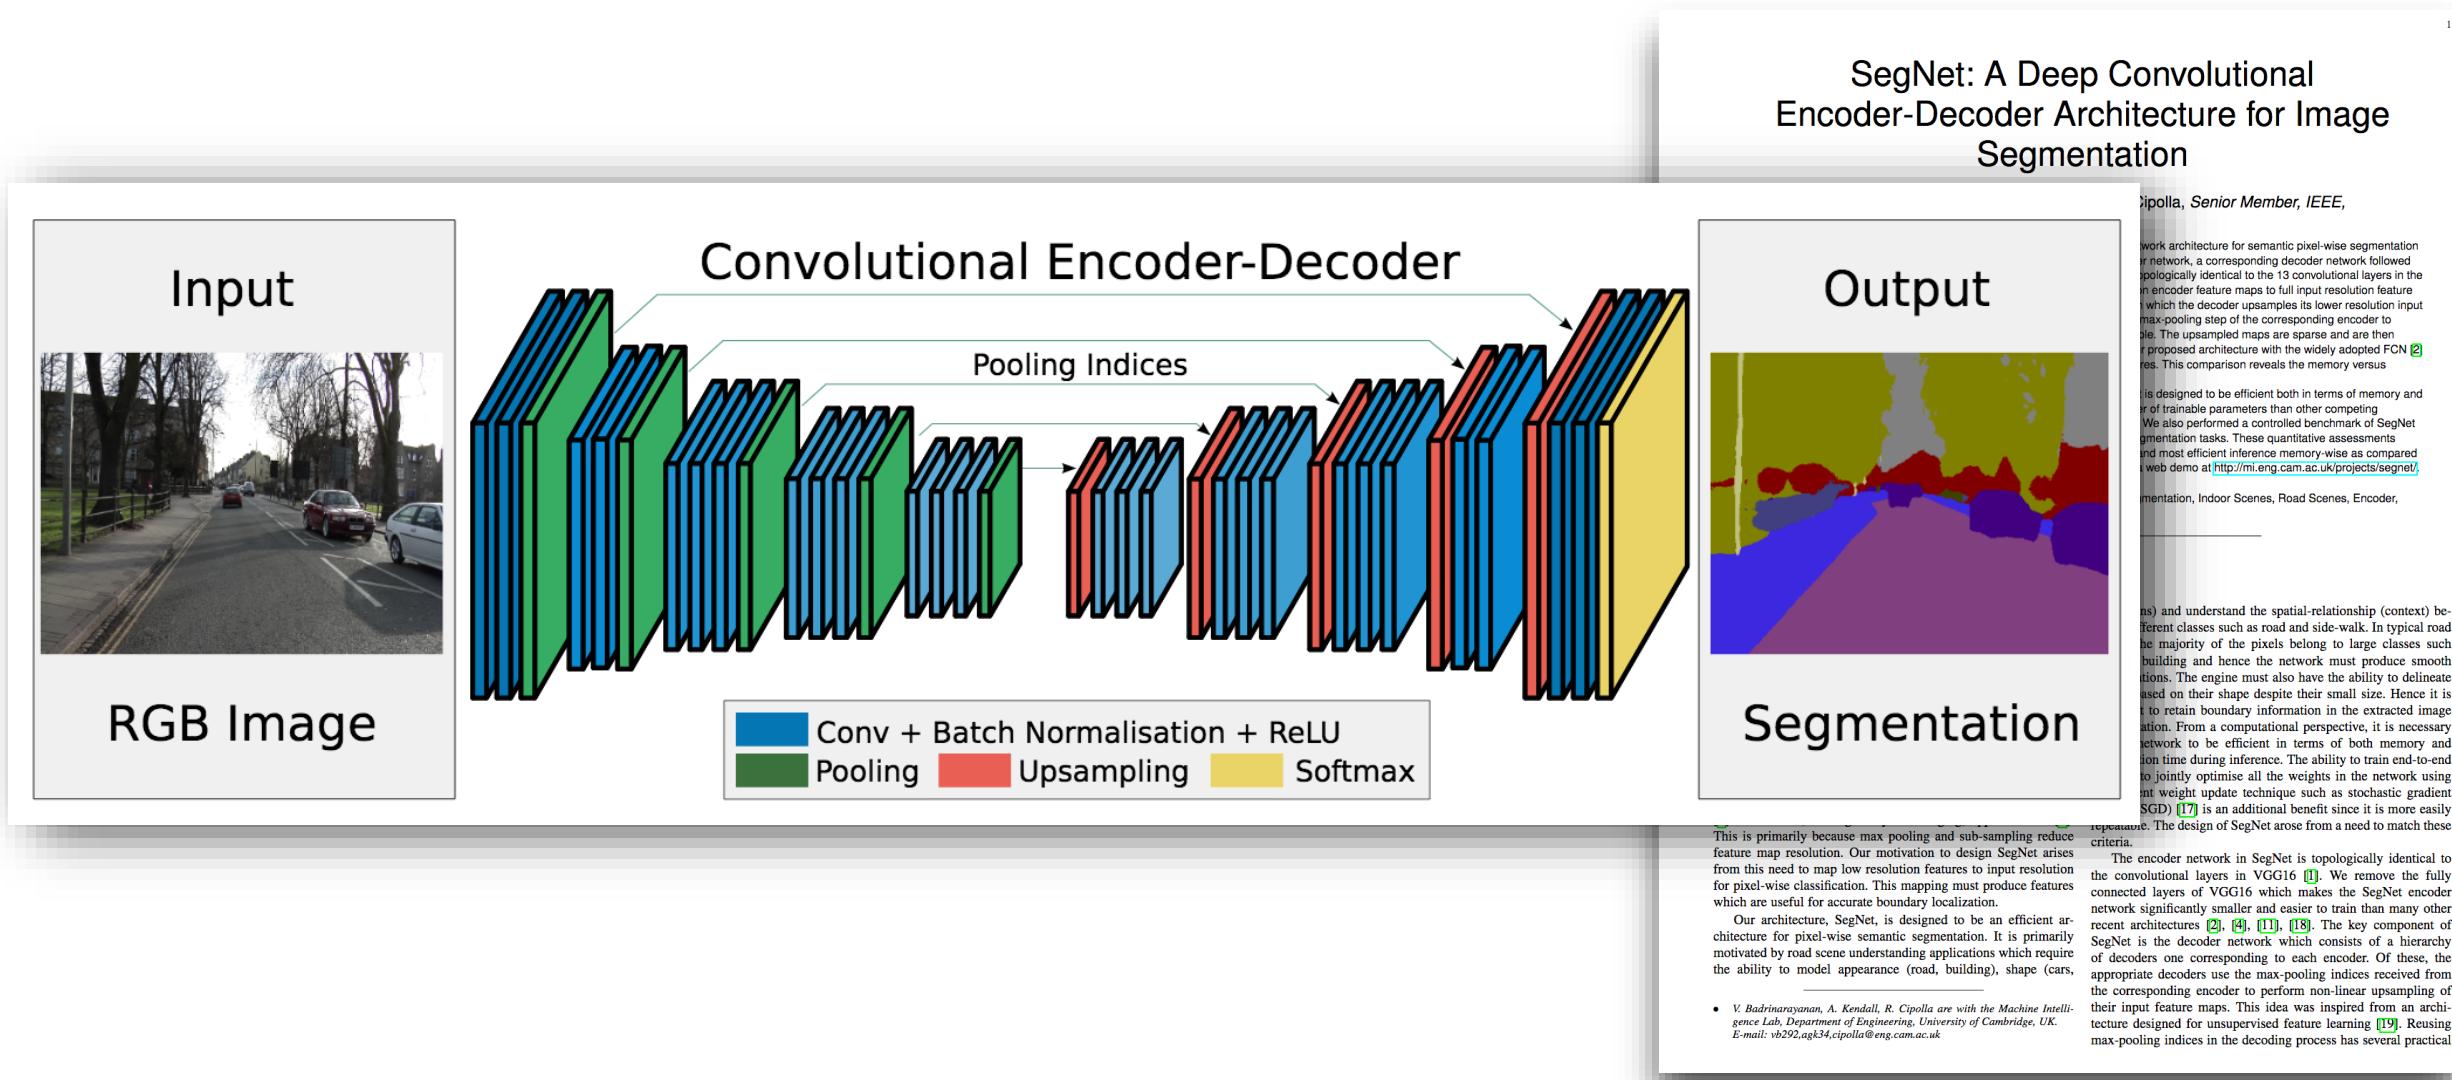
Recap: Transposed Convolution



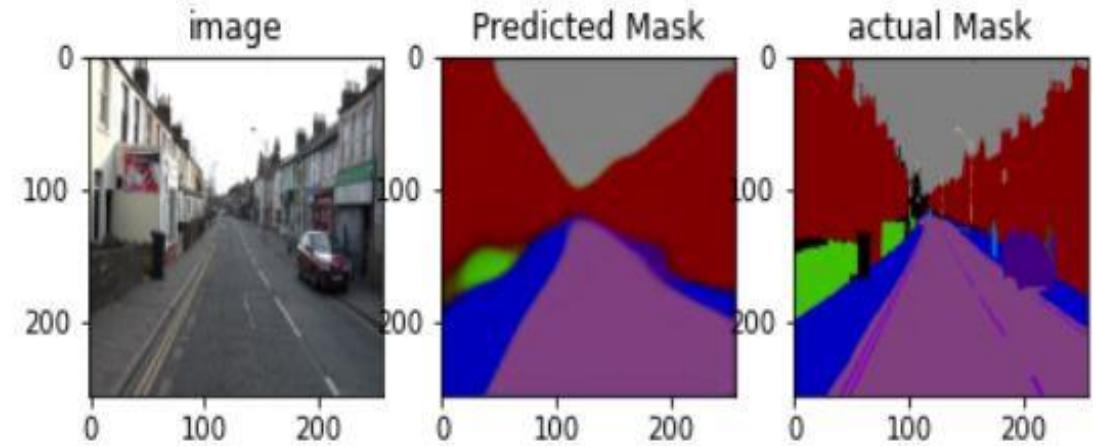
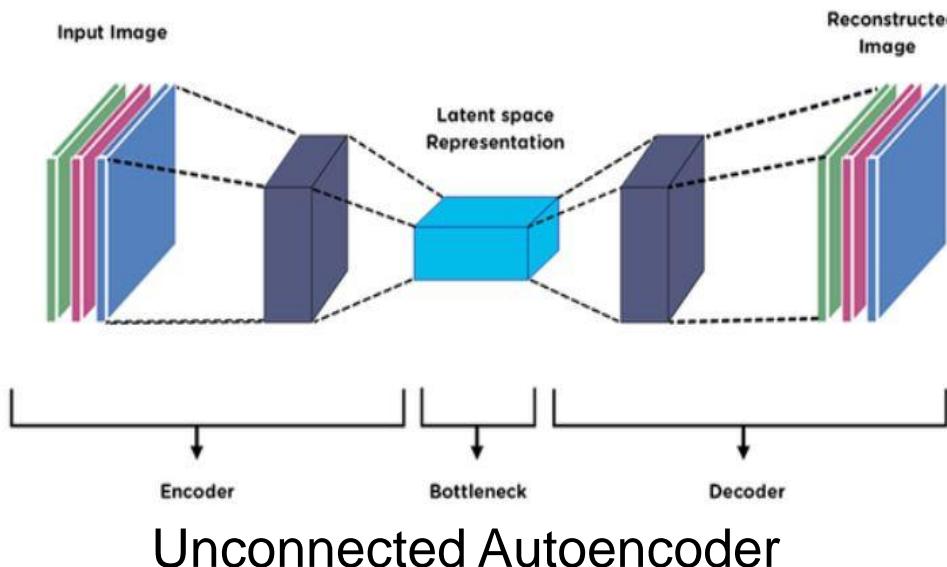
Transposed
Convolution used
for Upsampling in
CCNs



Autoencoder: Encoder-Decoder Architectures



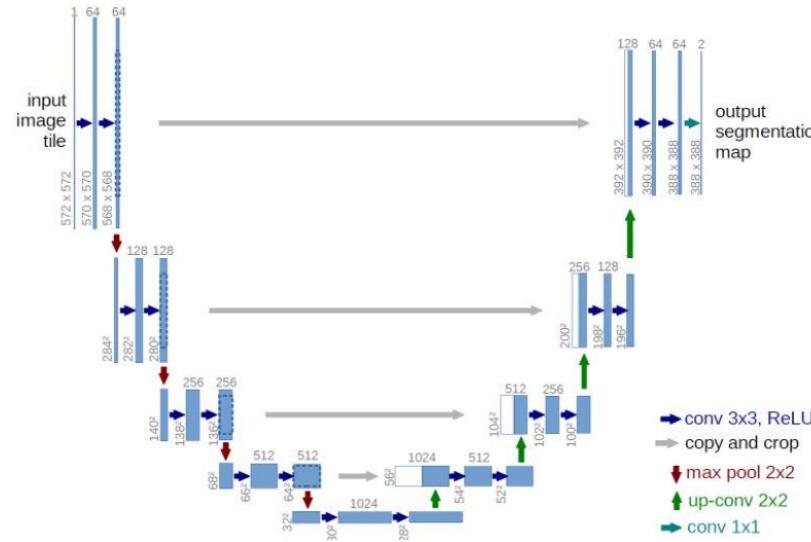
Unconnected Autoencoders



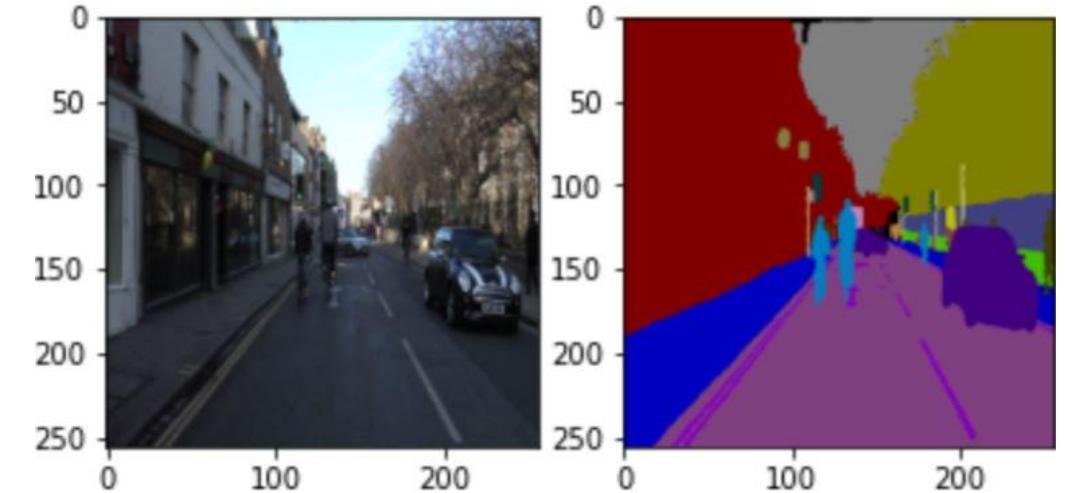
Imprecision in predicted Segmentation

However, convolution and pooling during encoding and transposed convolution during decoding leads to a imprecise feature prediction

Connected Autoencoders (U-Nets)



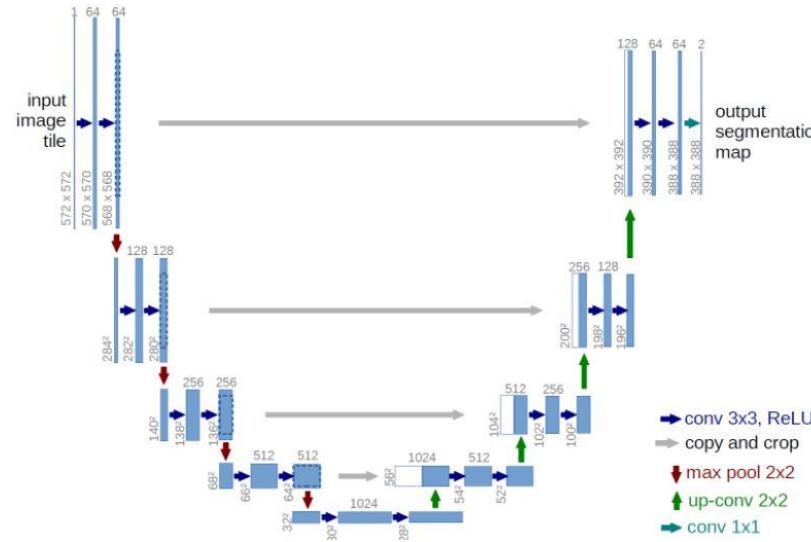
Connected Autoencoder (U-Net)



Predicted Segmentation

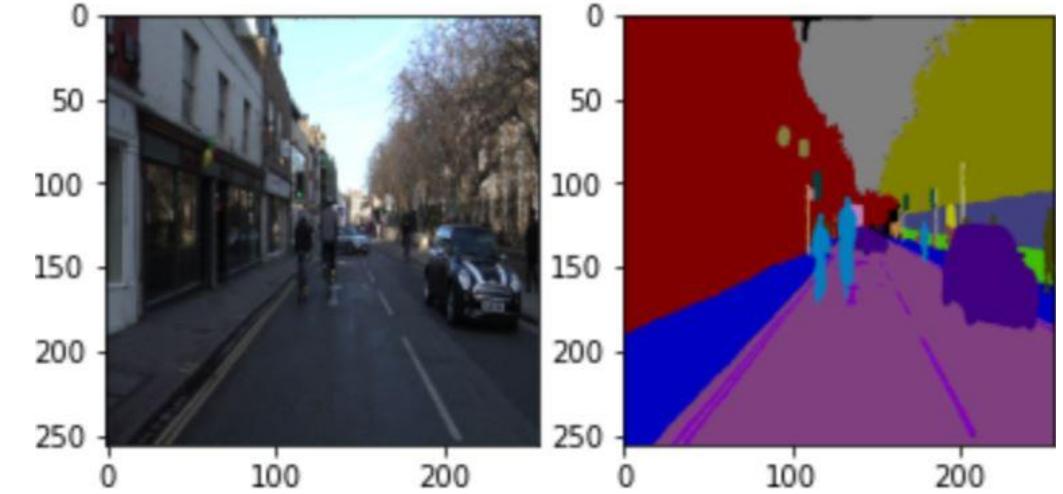
- U-Nets overcome this problem by connecting corresponding encoder-decoder layers with skip connections:
 - the output of an encoder level is skip-connected (concatenation) with the input of the corresponding decoder level

Connected Autoencoders (U-Nets)



Connected Autoencoder (U-Net)

- Why conv 1x1?
 - linear projection of stack of features (dimensionality reduction)



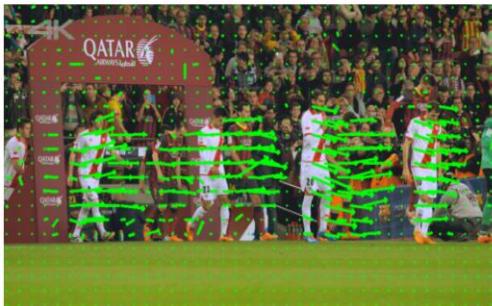
Predicted Segmentation

Course Overview

CW	Topic	Date	Place	Lab
41	Introduction and Course Overview	07.10.2025	Zoom	Lab 1
42	Capturing Digital Images	14.10.2025	Zoom	Lab 2
43	Digital Image Processing	21.10.2025	Zoom	Assignment 1
44	Machine Learning	28.10.2025	Zoom	
45	Feature Extraction	04.11.2025	Zoom	Open Lab 1
46	Segmentation	11.11.2025	Zoom	Assignment 2
47	Optical Flow	18.11.2025	Zoom	Open Lab 2
48	Object Detection	25.11.2025	Zoom	Assignment 3
49	Multi-View Geometry	02.12.2025	Zoom	Open Lab 3
50	3D Vision	09.12.2025	Zoom	Assignment 4
3	Trends in Computer Vision	13.01.2026	Zoom	
4	Q&A	20.01.2026	Zoom	Open Lab 4
5	Exam	27.01.2026	HS1 (Linz), S1/S3 (Vienna), S5 (Bregenz)	
9	Retry Exam	24.02.2026	tba	

Next Week: Optical Flow

What is Optical Flow?



Sparse Optical Flow (Flow Vector per Region)

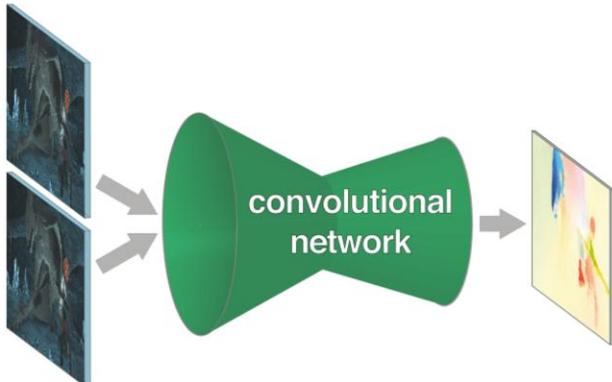


Dense Optical Flow (Flow Vector per Pixel)

Motion Vector of Pixel in Time Series (two consecutive Video Frames at Times t and t+1)

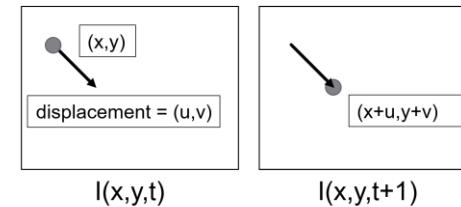
JKU JOHANNES KEPLER
UNIVERSITY LINZ

Optical Flow and Machine Learning



Encoder+Decoder Architectures (e.g. U-Nets)

The Optical Flow Equation



$I(x,y,t)$

$I(x,y,t+1)$

Brightness Constancy:

$$I(x + u, y + v, t + 1) = I(x, y, t)$$

$$0 \approx I(x + u, y + v, t + 1) - I(x, y, t)$$

Taylor Expansion:

$$\approx I(x, y, t + 1) + I_x u + I_y v - I(x, y, t)$$

$$= I(x, y, t + 1) - I(x, y, t) + I_x u + I_y v$$

Optical Flow Equation:

$$0 = I_t + I_x u + I_y v$$

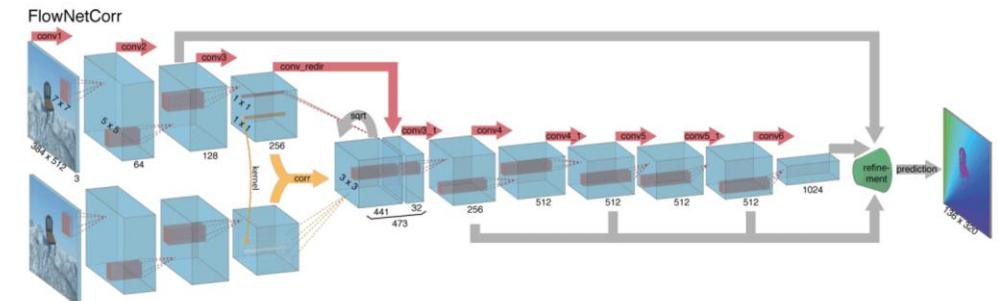
$$= I_t + \nabla I \cdot [u, v]$$

I_t, I_x, I_y are partial derivatives of image intensity (gradients) in t, x, y

JKU JOHANNES KEPLER
UNIVERSITY LINZ

Example: FlowNetCorr (Correlation)

<https://lmb.informatik.uni-freiburg.de/Publications/2015/DFIB15/flownet.pdf>



Input: 2 Tensors of individual RGB Images (Feature Maps are computed later → Correlation Layer)

JKU JOHANNES KEPLER
UNIVERSITY LINZ

Thank You

