

PROBABILISTIC MODELS – PART 5: LEARNING BAYESIAN NETWORKS

Gerhard Widmer

Institute of Computational Perception
Johannes Kepler University
Linz, Austria

gerhard.widmer@jku.at
www.cp.jku.at/people/widmer



October 28, 2025

Presentation partly based on and inspired by [Koller & Friedman, 2009] and [Russell & Norvig, 2021], including the use of some figures from their books and/or lecture slides.

Many thanks to Daphne Koller, Nir Friedman, Stuart Russell, and Peter Norvig for making these available
(pgm.stanford.edu; aima.cs.berkeley.edu).

Do not distribute!

Goals of this Lecture

- ▶ Discuss some fundamental issues related to learning from examples
- ▶ Introduce a criterion to be optimised: the Likelihood
- ▶ Explain the bias-variance trade-off and the need for regularisation

- ▶ Preview of next chapters:
 - Methods for learning the parameters (CPDs) of a model:  **Part 5a**
 - Methods for learning the structure of a model:  **Part 5b**

Outline

1 General Setting

Model Learning as an Optimisation Problem

A Possible Objective Function: The Likelihood

2 Overfitting, Generalisation, and the Bias-Variance Tradeoff

Overfitting vs. Generalisation

Bias vs. Variance

3 Preview of Learning Tasks

Motivation

Problem:

- ▶ “Manual” construction of model for a given problem may be impossible
 - ... because experts not available (or too expensive)
 - ... or problem not well enough understood or too complex

Possible Resource:

- ▶ May have a set of **example cases** (atomic events) observed in / collected from the world, e.g.:
- ▶ *Medical diagnosis*: database of patient records, patient histories, symptoms, tests performed, diagnoses, treatments, treatment outcomes, ...
- ▶ *Speech recognition*: recordings of speech annotated with word labels, sentence markers etc.

Goal:

- ▶ Construct a structured model of the (hidden) distribution most likely underlying the observed examples

 **Automatic Model Learning**

General Setting

Assumptions:

- ▶ World to be modelled is governed by some ‘true’, but unknown **distribution** P^* corresponding to some ‘true’, but unknown network $\mathcal{M}^* = (\mathcal{G}^*, \theta^*)$
- ▶ Given: a set $\mathcal{D} = \{x_1, \dots, x_M\}$ of M events coming from P^* (“training set”)
- ▶ Training instances x_i are all sampled, independently, from the same P^* : they are **independent and identically distributed (i.i.d.)**
- ▶ In words: The training examples are **representative** of the world P^*

Task:

- ▶ Learn some model $\tilde{\mathcal{M}}$ (from a given family of models) with a **distribution** $P_{\tilde{\mathcal{M}}}$ that is an approximation to P^* , and with a graph structure $\tilde{\mathcal{G}}$ that reflects the true **(in)dependencies** in the world.

May want to learn

- ▶ only model parameters $\tilde{\theta}$ for a fixed (given) structure, or both structure $\tilde{\mathcal{G}}$ and suitable parameters $\tilde{\theta}$
- ▶ a single model, or a whole spectrum (set) of different possible models (e.g., a probability distribution over models)

Learning as an Optimisation Problem

What model do we want to learn?

- ▶ Answer 1: “*The correct one*” — not decidable
- ▶ Answer 2: “*One that agrees well with the observed events \mathcal{D}* ”
- ▶ Note: The example observations \mathcal{D} is all the information we have ...

General approach:

- ▶ Define an “**objective function**” $F(\mathcal{M}, \mathcal{D})$ – a measure that estimates how ‘good’ a given model \mathcal{M} is in relation to the given training examples \mathcal{D}
- ▶ Develop an algorithm to find the model that maximises F :

$$\tilde{\mathcal{M}} = \arg \max_{\mathcal{M}} F(\mathcal{M}, \mathcal{D})$$

Learning is a **search / optimisation problem**:

- ▶ Search for a model $\tilde{\mathcal{M}}$ (in *huge* space of possible model structures and parameter settings) with maximum $F(\mathcal{M}, \mathcal{D})$

A Common Objective Function: The Likelihood

Definition

The **Likelihood** of a model \mathcal{M} relative to a dataset \mathcal{D} is **the probability that the model assigns to the set \mathcal{D}** :

$$L(\mathcal{M} : \mathcal{D}) = P_{\mathcal{M}}(\mathcal{D})$$

In Words: The Likelihood is ...

- ▶ the probability that one would get exactly the instances in \mathcal{D} if one randomly sampled $k = |\mathcal{D}|$ samples from the distribution $P_{\mathcal{M}}$
 - ▶ the probability that \mathcal{M} would generate exactly the observations \mathcal{D} if asked to generate k samples
 - ▶ the degree to which \mathcal{M} “fits” or “explains” the observations \mathcal{D} .
- 👉 **Search for model $\tilde{\mathcal{M}}$ with maximal likelihood relative to the given (fixed) training data \mathcal{D} .**

Calculating the Likelihood

Definition

If the examples $\mathcal{D} = \{x_1, \dots, x_M\}$ are **independent and identically distributed (i.i.d)**, the **likelihood** $L(\mathcal{M} : \mathcal{D})$ is

$$L(\mathcal{M} : \mathcal{D}) = P_{\mathcal{M}}(\mathcal{D}) = \prod_{x_i \in \mathcal{D}} P_{\mathcal{M}}(x_i)$$

In Words:

- ▶ The likelihood is the product of the probabilities assigned by the model to the individual training examples
- ▶ Easy to calculate for a given model and a training set.

The Log-Likelihood

Numerical problems with the likelihood function:

- ▶ Probability $P_{\mathcal{M}}(\mathcal{D})$ will be minuscule for any specific set of observations \mathcal{D}
- ▶ Trying to calculate this will produce an arithmetic underflow
- ▶ Solution: Use logarithm instead.

Definition

The **Log-Likelihood** $\ell(\mathcal{M} : \mathcal{D})$ of a model \mathcal{M} relative to a dataset \mathcal{D} is the logarithm of the likelihood:

$$\ell(\mathcal{M} : \mathcal{D}) = \log L(\mathcal{M} : \mathcal{D}) = \log \prod_{\mathbf{x}_i \in \mathcal{D}} P_{\mathcal{M}}(\mathbf{x}_i) = \sum_{\mathbf{x}_i \in \mathcal{D}} \log P_{\mathcal{M}}(\mathbf{x}_i)$$

Note:

- ▶ Likelihood and log-likelihood are monotonically related:
 $\ell(\mathcal{M} : \mathcal{D})$ has its maximum where $L(\mathcal{M} : \mathcal{D})$ is maximal
- 👉 **The log-likelihood will form a part of our objective function for learning structured probabilistic models.**

Overfitting vs. Generalisation

Consider a simple example:

- ▶ Want to learn a distribution P^* over a probability space defined by 20 binary variables
- ▶ There are $2^{20} > 10^6$ possible distinct atomic events in this space
- ▶ Training set \mathcal{D} consists of 1000 (different) instances.

Problem:

- ▶ If we permit our model class to contain *any distribution* possible over this space, the model \mathcal{M}_{ML} that maximises the (log-)likelihood would be one that
 - assigns equal probability 0.001 to each of the observed training instances
 - and probability 0 to all $(2^{20} - 1000)$ other events.

Resulting model \mathcal{M}_{ML} ...

- ▶ tightly fits and reproduces the training examples with high probability
- ▶ but assigns zero probability to (i.e., considers *impossible*) anything else.

☞ **Needed: A model that generalises!**

Overfitting vs. Generalisation

Generalisation:

- ▶ Given training set \mathcal{D} will usually not contain *all possible* situations (events) that could ever occur
- ▶ Purpose of a learned model is to answer queries about new situations
- ▶ Model must be more general than simple summary of training set!

Overfitting:

- ▶ Model with highest (log-)likelihood generally is one that exactly fits the data
- ▶ Assigns high probability to the seen examples, and low/zero probability to any other event
- ▶ Not useful for answering queries about new situations
- ▶ If we permit arbitrarily complex models and maximise the (log-)likelihood, we will get a model that overfits



Need some restrictions on allowed models!

Overfitting and Model Complexity

Note general relation between **Overfitting** and **Model Complexity**:

- ▶ Fitting a given training set \mathcal{D} with perfect precision requires complex models
(Model complexity = number of parameters required to specify the model)
- ▶ Overfitting models are usually complex models
- ▶ Reducing model complexity reduces ability of model to describe \mathcal{D} precisely
- ▶ Reducing model complexity enforces generalisation

For Bayesian Network Models:

- ▶ Complexity directly relates to number of edges in the graph
 - ▶ More parents per variable \Rightarrow more numbers in the CPDs
 - ▶ Reducing number of parents introduces stronger independencies
 - ▶ ... and thus makes it impossible to represent distributions where there are no independencies
- 👉 **Constraining the structure of a BN reduces the space of representable distributions**

Re: “Overfitting Models are Complex”

Consider a simple example:

- ▶ World with three binary variables A, B, C
- ▶ Given training set \mathcal{D} :

	A	B	C
x_1	1	0	0
x_2	0	1	0
x_3	0	0	1

Exercise:

- ▶ Construct a Bayesian network over A, B, C that assigns probability $1/3$ to each of the three examples x_1, x_2, x_3 and 0.0 to any other event $\in \{0, 1\}^3$
- ▶ Can this be done without creating a fully connected network?

Bias and Variance

Goal:

- ▶ Avoid learning an overfitting model
- ▶ Force learner to generalise

Approach:

- ▶ Put constraints on class of models allowed to the learner:
- ▶ **Hard constraints:** Strictly restrict the class of models (e.g., only permit certain structures, limit number of parents, etc.)
- ▶ **Soft constraints:** Introduce an additional *regularisation term* to the objective function that adds a penalty for complex models.

Consequence:

- ▶ Models from a constrained model class are limited in how closely they can approximate the target distribution P^*
- ▶ Learned models will not be able to precisely describe the empirical distribution in the training data, and perhaps also the target distribution P^*

Error possibility introduced by restricting expressivity of model class is called **Bias.**

Bias and Variance

On the other hand:

- ▶ A large space of highly expressive (complex) models is more likely to contain a model \tilde{M} that closely approximates P^* (low bias)

But:

- ▶ Limited training set \mathcal{D} may not be able to select the ‘right’ model among the large number of models in the hypothesis space
- ▶ Many candidate models will have similar likelihood relative to \mathcal{D}
- ▶ Small changes in \mathcal{D} can radically change the properties of the selected model
- ▶ Running the learning algorithm several times with different training sets sampled from P^* will produce highly variable overfitting models.

Error possibility introduced by permitting high expressivity of model class is called **Variance.**

The Bias-Variance Tradeoff

Fundamental Tradeoff:

- ▶ A restriction to **simple models** makes hypothesis space smaller and increases the probability of **bias error**
- ▶ On the other hand, in a smaller hypothesis space, it is less likely to find an overfitting model.

vs.

- ▶ Permitting **complex models** reduces probability of bias error
- ▶ but introduces **variance** as a potential source of error:
- ▶ There may be many models that fit \mathcal{D} “*by chance*” ...

- ⇒ **Tradeoff:** generalisation vs. overfitting, bias vs. variance
- ⇒ Task of system designer/experimenter:
find a good place in this continuum, relative to given problem (data set)
- ⇒ There is no theoretical decision guide
- ⇒ Much of Machine Learning is about this ...

Preview: Learning Tasks in Graphical Models

Different learning scenarios:

- ▶ **Parameter Learning** (given a fixed model structure)
 - ▶ Maximum likelihood parameter estimation
 - ▶ Bayesian parameter estimation
- ▶ **Structure Learning**
(learning both the structure and the parameters of a model)

Additional level of difficulty: **Learning from Incomplete Observations**

In this class:

- ▶ Parameter learning in Bayesian networks (☞ Part 5.a)
- ▶ Structure learning in Bayesian networks (☞ Part 5.b)
- ▶ Learning from incomplete observations in a special class of graphical models: Hidden Markov Models (HMMs) (☞ Part 6.a)

What you should remember of this section

- ▶ The General Setting: Model Learning as an Optimisation Problem
- ▶ Objective Functions: Likelihood and Log-likelihood
- ▶ Generalisation vs. Overfitting
- ▶ Bias vs. Variance

Literature

Koller, Daphne and Friedman, Nir (2009).
Probabilistic Graphical Models: Principles and Techniques. Cambridge, MA: MIT Press.